

Are Relational Databases Obsolete?



Rudi Bruchez

SQL SERVER CONSULTANT AND TRAINER

www.linkedin.com/in/rudibruchez/ www.babaluga.com



Summary



What about NoSQL?

Are SQL databases obsolete?

SQL database are not web scale

Joins are slow


The relational schema is too rigid

What is the CAP theorem?

SQL is very old




Can We Predict the Load?



Traditional application
You can predict your
load




Web scale
On the web or mobile,
you can't predict the
load




Elasticity
You need to be able to
scale




Do We Need Scalability?



Is your data big?
Maybe you don't need
Big Data



Complexity
Scaling means
creating more
complexity



NoSQL and Big Data
They are good at
scaling



What Is NoSQL?




It is a movement




Key-value store




Document database



Wide-row store




Graph database




Search engine



Is It Really No SQL?



No to SQL?
We don't like this
language



Not only SQL?
We get more choice



But they have SQL!
So we love this
language



Apache Cassandra's CQL

```
CREATE TABLE enrollment
(
    contactid uuid,
    enrollment_date timestamp,
    session text,
    PRIMARY KEY (contactid, enrollment_date)
) WITH CLUSTERING ORDER BY (enrollment_date DESC);

INSERT INTO enrollment (contactid, enrollment_date, session)
VALUES (
    uuid(),
    toTimestamp(now()),
    'Cassandra Fundamentals'
);

SELECT *
FROM enrollment
WHERE contactid = 66fedeb0-69db-41df-b10f-cec2d42a77bd;
```



Couchbase Server's N1QL

```
SELECT DISTINCT(name) AS Airline, id
FROM `travel-sample`
WHERE type = 'airline'
AND country LIKE '%Fran%' LIMIT 5 OFFSET 5
```

```
SELECT a.name, s.flight, s.utc, r.sourceairport,
       r.destinationairport, r.equipment
FROM `travel-sample` r UNNEST r.schedule s
JOIN `travel-sample` a ON KEYS r.airlineid
WHERE r.sourceairport='SEA'
AND r.destinationairport='MCO' AND s.day=6
ORDER BY a.name
```



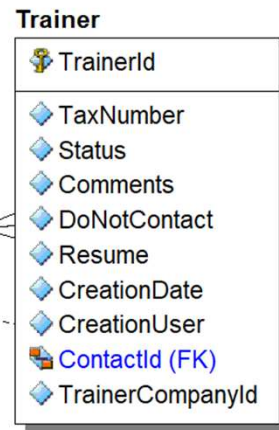
It Is Not About the Language

SQL is well-known, we can build on that
In Big Data, Hive or Spark implement SQL
NoSQL really means NoRelational
Different data models
Martin Fowler

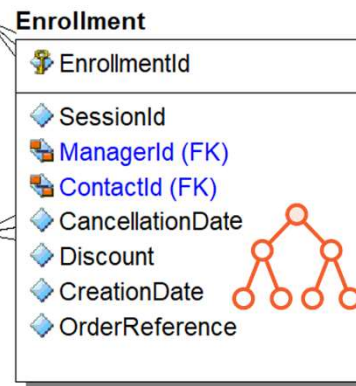

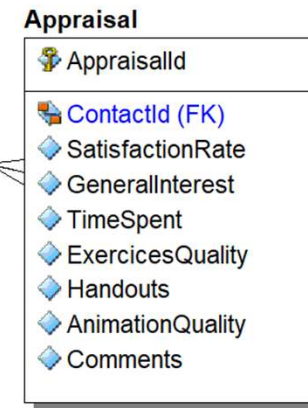


The Pachadata Model

```
SELECT c.ContactId
, c.Title
, c.LastName
, c.FirstName
, c.Email
, a.Address1
, c1.Name AS City
FROM Contact.Contact c
JOIN Contact.Address a
ON c.AddressId = a.AddressId
JOIN Reference.City c1
ON a.CityId = c1.CityId
```




```
SELECT *
FROM Trainer.Trainer t
JOIN Contact.Contact c
ON t.ContactId = c.ContactId
```




```
SELECT *
FROM Enrollment.Enrollment e
JOIN Contact.Contact c ON e.ContactId = c.ContactId
JOIN Course.Session s ON e.SessionId = s.SessionId
```




NoSQL Data Model



Faster and Simpler
Handle the load




Key-value store
Great performances



No constraint
Focus on the data



Relational Tables



A Document Database

```
{
  "_Id":8608,
  "Title":"Mr.",
  "LastName":"Dauriac",
  "FirstName":"Alfred",
  "Email":"a.gauthier@cookingschool.com",
  "Phone":"+625166031",
  "Gender":"H",
  "Mobile":"+994997445",
  "Address": {
    "Address1":"Common Street 33",
    "CityName":"San Gommery",
    "ZipCode":"56920"
  },
  "Enrollments": [
    {
      "SessionDate": "2018-07-05",
      "Language":"EN",
      "Class":38,
      "Price":2500.00,
      "Titre":"MySQL 5.7 essential"
    }
  ]
}
```

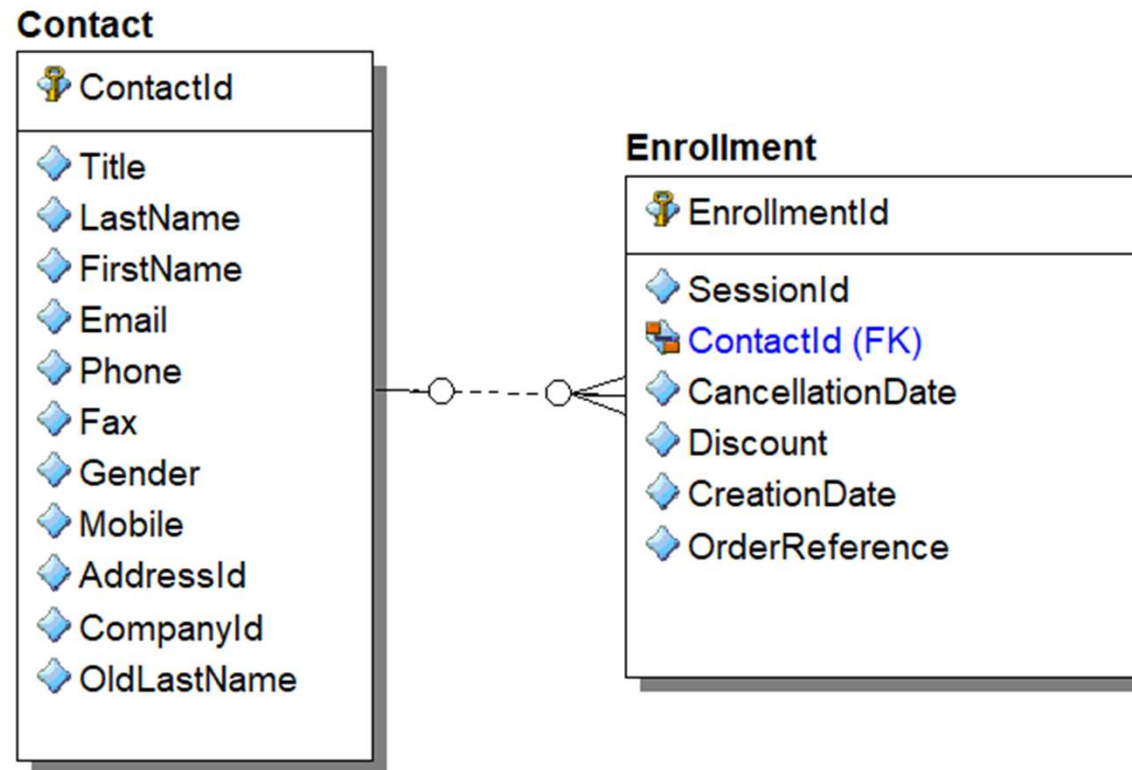


A Document Database

```
{
  "_Id":8608,
  "Title":"Mr.",
  "LastName":"Dauriac",
  "FirstName":"Alfred",
  "Email":"a.gauthier@cookingschool.com",
  "Phone":"+625166031",
  "Gender":"H",
  "Mobile":"+994997445",
  "Address": {
    "Address1":"Common Street 33",
    "CityName":"San Gommery",
    "ZipCode":"56920"
  },
  "Enrollments": [
    {
      "SessionDate": "2018-07-05",
      "Language":"EN",
      "Class":38,
      "Price":2500.00,
      "Titre":"MySQL 5.7 essential"
    },
    {
      "SessionDate": "2018-11-22",
      "Language":"EN",
      "Class":15,
      "Price":2500.00,
      "Titre":"MySQL 5.7 advanced"
    }
  ]
}
```



Relational Tables



Domain-Driven Design Fundamentals

by Steve Smith and Julie Lerman

This course teaches the fundamentals of Domain-Driven Design (DDD) through a demonstration of customer interactions and a complex demo application, along with advice from Eric Evans.

 Start Course



Bookmark



Add to Channel



Live mentoring

Table of contents

Description

Transcript

Exercise files

Discussion

Learning Check

Recommended

This course is part of:  ASP.NET MVC 5 Path

[Expand all](#)



Introducing DDD



24m 19s



DDD: Modeling Problems in Software



45m 6s



Course authors



Steve Smith



Julie Lerman

Course info

Level Intermediate

Rating ★★★★★ (1396)

My rating ★★★★★

Duration 4h 16m


Released 25 Jun 2014

Share course

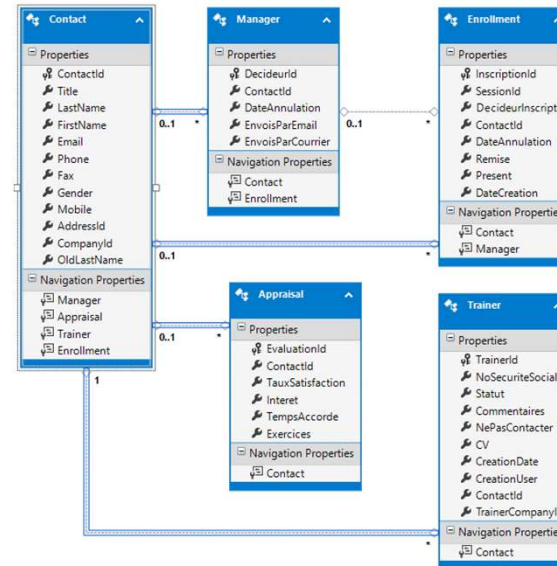


Domain Driven Design Aggregates

Relational Database

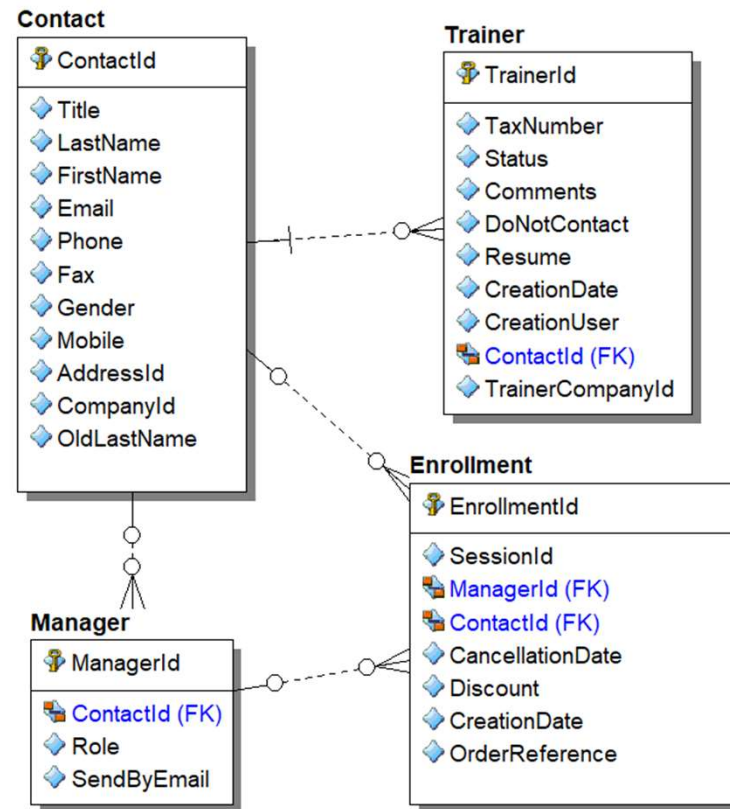



Entity Framework



Joins as Hierarchy

```
SELECT c.Title
      ,c.LastName
      ,c.FirstName
      ,c.Email
      ,s.CourseId
      ,s.LanguageCd
      ,s.RoomId
      ,s.StartDate
FROM Contact.Contact c
JOIN Enrollment.Enrollment e
      ON c.ContactId = e.ContactId
JOIN Course.Session s
      ON e.SessionId = s.SessionId
WHERE c.ContactId = 8608
```

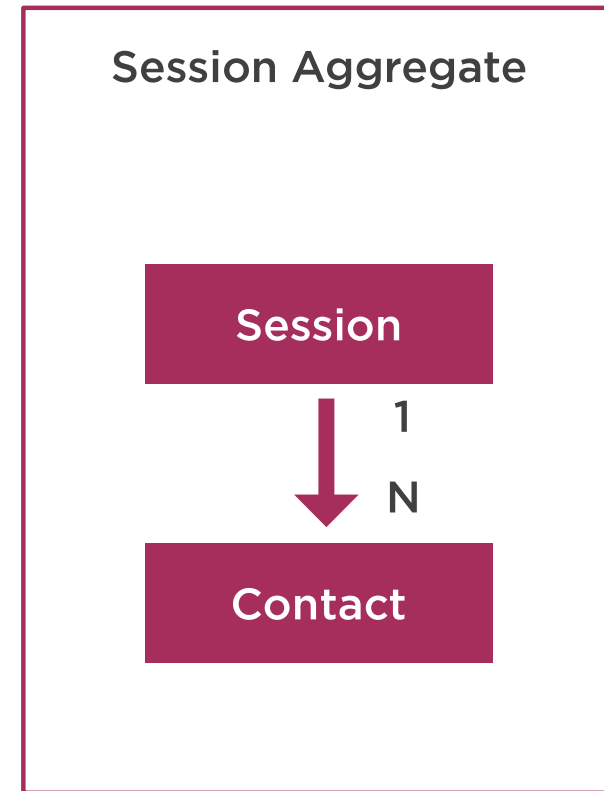
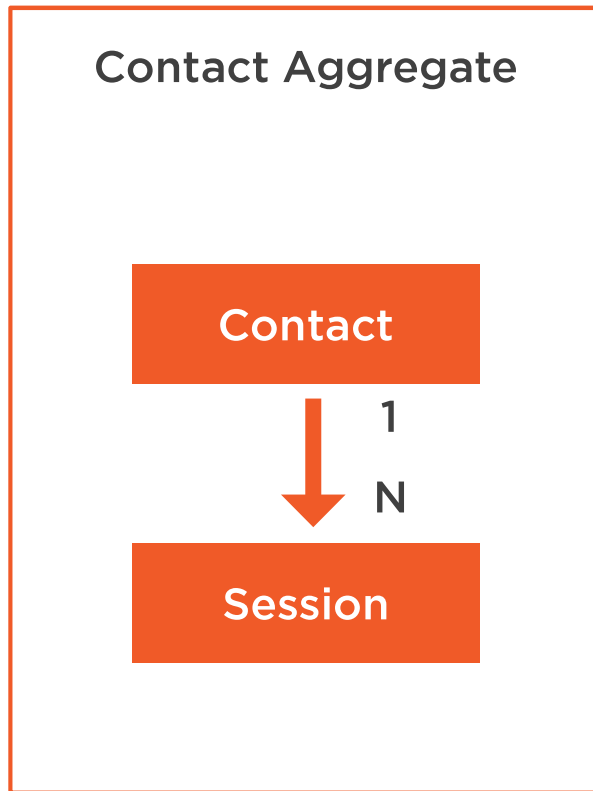


Domain Driven Design Aggregates

Reduce the complexity of a data model
Object hierarchy
Bounded context for the client application
Aggregate root

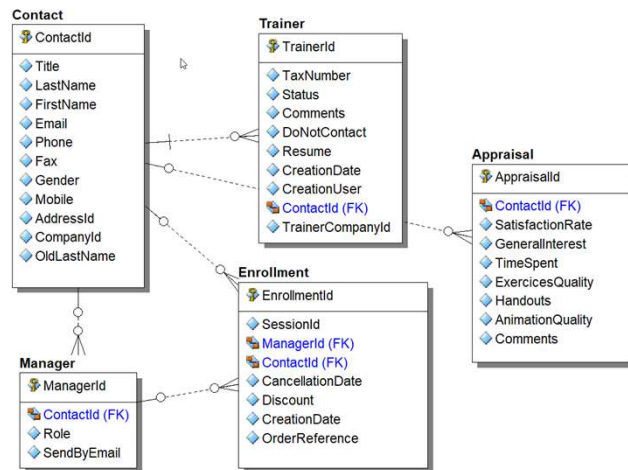


Pachadata Aggregates




Domain Driven Design Aggregates

Relational Database



SELECT
FROM
JOIN

Session Aggregate



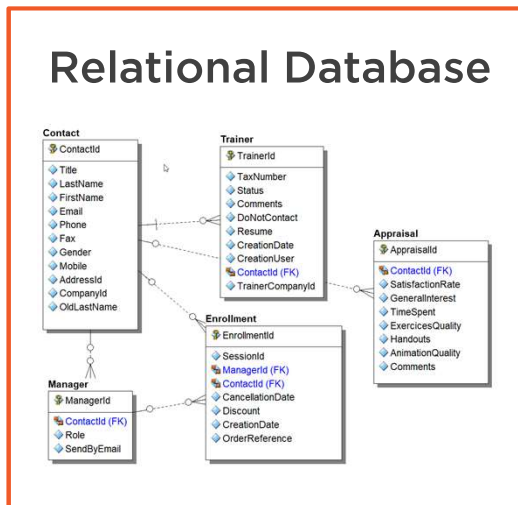
Relational Databases as Integration Stores

Contact
Management

Session
Management

Trainer
Management

Enrollment
Management



NoSQL Databases as Aggregates

Contact
Management

Session
Management

NoSQL Database

```
{
  "_id": 8608,
  "Title": "Mr.",
  "LastName": "Dauriac",
  "FirstName": "Alfred",
  "Email": "a.gauthier@cookingschool.com",
  "Phone": "+625166031",
  "Gender": "H",
  "Mobile": "+994997445",
  "Address": {
    "Address1": "Common Street 33",
    "CityName": "San Gommery",
    "ZipCode": "56920"
  },
  "Enrollments": [
    {
      "SessionDate": "2018-07-05",
      "Language": "EN",
      "Class": 38,
      "Price": 2500.00,
      "Titre": "MySQL 5.7 essential"
    }
  ]
}
```

Trainer
Management

Enrollment
Management



NoSQL Databases as Aggregates

Contact
Management



Session
Management

NoSQL Database

```
{
  "_id": 8608,
  "Title": "Mr.",
  "LastName": "Dauriac",
  "FirstName": "Alfred",
  "Email": "a.gauthier@cookingschool.com",
  "Phone": "+625166031",
  "Gender": "H",
  "Mobile": "+994997445",
  "Address": {
    "Address1": "Common Street 33",
    "CityName": "San Gommery",
    "ZipCode": "56920"
  },
  "Enrollments": [
    {
      "SessionDate": "2018-07-05",
      "Language": "EN",
      "Class": 38,
      "Price": 2500.00,
      "Titre": "MySQL 5.7 essential"
    }
  ]
}
```




NoSQL Database


```
{
  "_id": 44,
  "Title": "Mr.",
  "SessionDate": "2018-07-05",
  "Language": "EN",
  "Class": 38,
  "Price": 2500.00,
  "Titre": "MySQL 5.7 essential",
  "Enrollments": [
    {
      "LastName": "Dauriac",
      "FirstName": "Alfred",
      "Email": "a.gauthier@cookingschool.com",
      "Phone": "+625166031",
      "Gender": "H",
      "Mobile": "+994997445",
      "Address": {
        "Address1": "Common Street 33",
        "CityName": "San Gommery",
        "ZipCode": "56920"
      }
    }
  ]
}
```



Choosing NoSQL?



Flexibility
With different
databases




Complexity
More services, more
components



Problems
Get used to it




Choosing NoSQL?



Aggregates

You loose centrality



Scalability

Do you really need it?




Ease of use


Is it a valid argument?




Relational Databases Are Too Rigid



Flexibility
With schemaless
databases



With strong coupling?
All tiers need to
change



Schema-later
There is always a
schema



The Contact Document

```
{  
  "_Id":8608,  
  "Title":"Mr.",  
  "LastName":"Dauriac",  
  "FirstName":"Alfred",  
  "Email":"a.gauthier@cookingschool.com",  
  "Phone":"+625166031",  
  "Mobile":"+994997445"  
}
```



The Contact Document

```
{  
  "_Id" : 8608,  
  "Title" : "Mr.",  
  "LastName" : "Dauriac",  
  "FirstName" : "Alfred",  
  "Email" : "a.gauthier@cookingschool.com",  
  "Phone" : "+625166031",  
  "Mobile" : "+994997445",  
  "Gender" : "M"  
}
```



The Contact Document


```
var client = new MongoClient("mongodb://192.168.1.15:27017");
var database = client.GetDatabase("pachadata");
var collection = database.GetCollection<BsonDocument>("contacts");
var document = collection.Find(new BsonDocument()).FirstOrDefault();
BsonString gender;
if (!document.TryGetElement("Gender", out gender))
{
    // ...
}
```




What Do We Need?



No coupling
Tiers should not be too dependent



Database is a service
Decoupling database from client code



Abstraction layer
Between the database and the client code



Using Views

```
CREATE VIEW Contact.ContactAggregate  
AS
```

```
    SELECT c.Title, c.LastName,  
           c.FirstName, c.Email,  
           s.CourseId, s.LanguageCd,  
           s.RoomId, s.StartDate
```

```
FROM Contact.Contact c
```

```
LEFT JOIN Enrollment.Enrollment e  
    ON c.ContactId = e.ContactId
```

```
LEFT JOIN Course.Session s  
    ON e.SessionId = s.SessionId;
```

```
CREATE VIEW Course.SessionAggregate  
AS
```

```
    SELECT s.SessionId, s.CourseId,  
           s.LanguageCd, s.RoomId,  
           s.StartDate, c.Title,  
           c.LastName, c.FirstName,  
           c.Email
```

```
FROM Course.Session s
```

```
LEFT JOIN Enrollment.Enrollment e  
    ON e.SessionId = s.SessionId
```

```
LEFT JOIN Contact.Contact c  
    ON c.ContactId = e.ContactId;
```



“All views that are theoretically updatable are also updatable by the system”

Edgar Codd

Rule 6 – The *view updating rule*



Conclusion




Should we always
follow trends?



Choose wisely

{JSON}

NoSQL movement



We keep SQL Server
for now



Choice
polyglot storage



Predictable
performances?



Colors

