

# Understanding Filtered Indexes

---



**Kimberly L. Tripp**

OWNER/PRESIDENT - SQLSKILLS.COM

@kimberlyltrippp

[www.sqlskills.com/blogs/kimberly](http://www.sqlskills.com/blogs/kimberly)



# Module Overview



**Filtered indexes and requirements**


**Filtered index use cases**

**Filtered indexes and interval  
subsumption**

**Filtered indexes and plan caching**

**Filtered index strategies**

# What is a Filtered Index?



Traditionally, row-based index has one entry for every table row (in index order) in the leaf level

The B-tree is used to navigate to the leaf level

What if you're always using an index for just one portion of data

- WHERE Active = 1
- WHERE Status > 8
- WHERE SalesPerson IS NOT NULL
- WHERE OrderDate BETWEEN ...

The point: covering ALL of the data is expensive!



# Filtered Indexes

Only applies to  
nonclustered  
indexes

Overall size might  
be significantly  
lower than an  
unfiltered index

Require consistent  
session settings at  
many levels of  
creation and use

Require code changes  
for some predicates  
and to use in stored  
procs / sp\_executesql

Maintenance costs  
lower as only DML that  
affects rows in index  
causes index changes

DTA can suggest  
filtered indexes,  
missing index DMVs  
do not



# Session Settings

Session settings control behavior, and the result of some computations

Data in these persisted structures must be consistent

Session settings that must be on:

- ANSI\_NULLS
- ANSI\_WARNING
- QUOTED\_IDENTIFIER
- CONCAT\_NULL\_YIELDS\_NULL
- ANSI\_PADDING
- ARITHABORT

## Msg 1934, Level 16, State 1, Line 1

CREATE INDEX failed because the following SET options have incorrect settings: 'QUOTED\_IDENTIFIER'. Verify that SET options are correct for use with indexed views and/or indexes on computed columns and/or filtered indexes and/or query notifications and/or XML data type methods and/or spatial index operations.

Session setting that must be off:

- NUMERIC\_ROUNDABORT



# Client Consistency

## Consistency at creation

Table and index must be created with correct session options

## Consistency with table changes

All DML statements must use correct session settings

## Consistency with query access

All SELECT statements must use correct session settings



# Use Case: Imagine “Types” of Data

For example, insurance types, document types, customer types

Imagine there are 9 “statuses” of employee records (distribution can be even or NOT – doesn’t matter)

- When you look at status = 1, you’re only interested in c6, c8, c4, c7
- When you look at status = 2, you’re only interested in c12, c16, c14, c6
- When you look at status = 3, you’re only interested in c2, c3, c4, c6

**FILTER** each index on status with **ONLY** relevant included columns for a much smaller index size (matching only number of rows for filter)

Each index is only  $1/n$  of the table; even with **n** specialized indexes storage is only what one unfiltered index would have been!



# Demo



**Filtered indexes: EmployeeCaseStudy**





# Cover More!

They're small so consider adding more columns to support additional queries!

More effective for covering when sets are clearly defined


If equality-based sets consider NOT including that column (unless query needs it)

If NOT equality-based sets consider including column as part of key for effective seeking




# Why Not One Per Set Always?

Leaf level of a regular nonclustered on  
(SalesPerson) INCLUDE (c6, c8)



Leaf level of individual filtered  
indexes by SalesPerson  
INCLUDE (C6, c8)




Size difference negligible (but don't need to include filter column)

- Predicate uses are **MORE** limited
- Covering choices are **WAY MORE** flexible

# When Is One Per Set Acceptable?

Leaf level of a regular nonclustered on  
(SalesPerson) INCLUDE (c6, c8)  
requires same scolumns for all  
salespeople



What if you want different  
columns per set?

- SalesPerson = 1 (c16, c22)
- SalesPerson = 2 (c2, c8)
- Etc...



# Demo



## Filtered indexes and interval subsumption



# Plan Caching

Stored procedures and `sp_executesql` statements compile and save an optimized plan for execution in the cache

- Not a permanent structure
- Plan is based on parameters used at time of plan caching (when there isn't already a plan in cache for that object)

Depending on how the criteria needing filtered index is specified, may or may not be able to leverage a filtered index without changing code

- Literals CAN use a filtered index
- Parameters CANNOT use filtered index unless `OPTION (RECOMPILE)` specified
- Variables CANNOT use filtered index unless `OPTION (RECOMPILE)` specified



# Demo



## Filtered indexes and plan caching



# What Filter?

sys.indexes

## Has two interesting columns:

- has\_filter: if index has a filter predicate
- filter\_definition: expression for filter definition



# What Filter?

sp\_helpindex

**Doesn't show included columns or filtered indexes**

- Use my tweaked “sp\_helpindex” to get better information and determine if one index really is redundant/duplicate:
- [https://www.sqlskills.com/BLOGS/KIMBERLY/category/sp\\_helpindex-rewrites](https://www.sqlskills.com/BLOGS/KIMBERLY/category/sp_helpindex-rewrites)





# What Filter?

DBCC  
SHOW\_STATISTICS

**Shows filter expression that defines subset over which statistics are computed**

- More accurate over the filtered set



# Demo



## Examining the filter predicate



# Maintenance

Statistics more accurate (when created) as they describe fewer rows

**More accurate after created or updated**

**Not auto-updated until threshold is met**

- Versions prior to SQL Server 2016, it's minimum of 500 + 20% of the TABLE
- Using trace flag 2371 OR in SQL Server 2016, it's "dynamic", happens much earlier than 20% for tables >25K rows

**Do not rely on "auto update statistics"**

**Better to have automated maintenance to keep statistics up-to-date**



# Maintenance

Index maintenance is  
required

**ALTER INDEX ... REBUILD/REORGANIZE  
work with filtered indexes**

**Online index operations work with filtered  
indexes**



Don't go wild with filtered indexes – powerful but definitely specific uses!

Test, test, test!



# What We Covered



**Filtered indexes and requirements**

**Filtered index use cases**

**Filtered indexes and interval  
subsumption**

**Filtered indexes and plan caching**

**Filtered index strategies**