

# Understanding the Behavior of the Query Optimiser

---



**Gail Shaw**

TECHNICAL LEAD

@SQLintheWild <http://sqlinthewild.co.za>



# Agenda



Overview of query execution process

How does the optimizer choose a plan?


Role of row estimations

Plan caching and reuse

Cause of bad parameter sniffing



# Query Execution Process



# Query Optimization

**Cost-based optimizer**

**Goal to find the lowest-cost plan in the time available**

**Costs based on row estimations**

**Row estimations from column statistics**



# How row estimations are calculated

RANGE_HI_KEY	RANGE_ROWS	EQ_ROWS	DISTINCT_RANGE_ROWS	AVG_RANGE_ROWS
CA	0	254235	0	1
EU	0	125516	0	1
MU	0	41	0	1
UK	0	254317	0	1
US	0	315891	0	1

## Clustered Index Scan (Clustered)

Scanning a clustered index, entirely or only a range.

Physical Operation	Clustered Index Scan
Logical Operation	Clustered Index Scan
Actual Execution Mode	Row
Estimated Execution Mode	Row
Storage	RowStore
Actual Number of Rows	315891
Actual Number of Batches	0
Estimated I/O Cost	6.68016
Estimated Operator Cost	6.94145 (14%)
Estimated Subtree Cost	6.94145
Estimated CPU Cost	0.261289
Estimated Number of Executions	1
Number of Executions	8
Estimated Number of Rows	315891
Estimated Row Size	28 B
Actual Rebinds	0
Actual Rewinds	0
Ordered	False
Node ID	9

## Predicate

[Nile].[dbo].[Orders].[Country] as [o].[Country]=[@Country]



# Demo



How row estimations change based on the parameter/constant/variable value



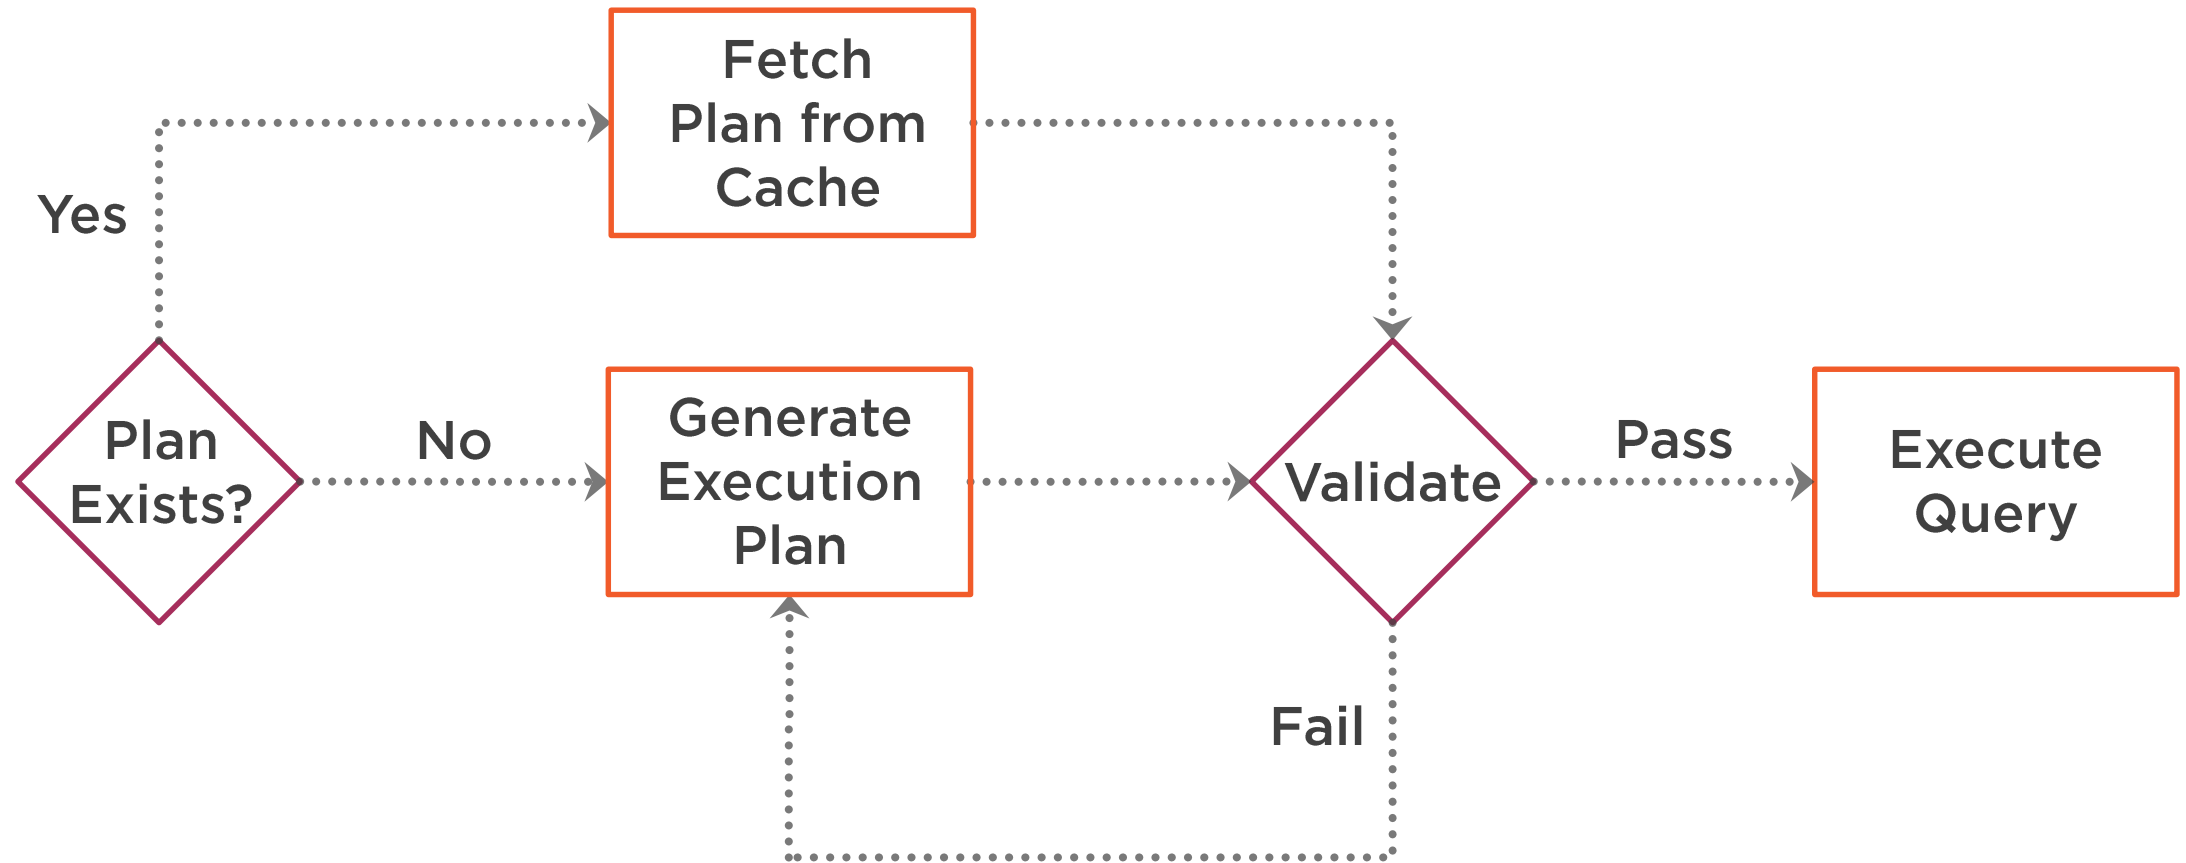
# Plan Caching

**Query Optimization is expensive**

**SQL Server caches execution plans to avoid unnecessary optimizations**



# Plan caching and reuse





# Demo



**Plan compilation, caching and reuse**



Bad parameter sniffing is caused by inappropriate reuse of cached plans.



# Summary



High-level look at query optimization  
How row estimations affect plan choices  
Caching and reuse of query plans

