

Row-based vs. Column-based Indexes



Kimberly L. Tripp

OWNER/PRESIDENT - SQLSKILLS.COM

@kimberlyltrippp

www.sqlskills.com/blogs/kimberly



Module Overview



Row-based index structures

Column-based index structures

Understanding your workload

Row-based vs. column-based indexes

Columnstore indexes by version


Indexed views vs. columnstore indexes




Table Structures



Heap structure
(possible in all versions)




**Row-based
clustered index**
(possible in all versions)



**Column-based
clustered index**
(limited options in SQL Server 2014;
more options in SQL Server 2016+)

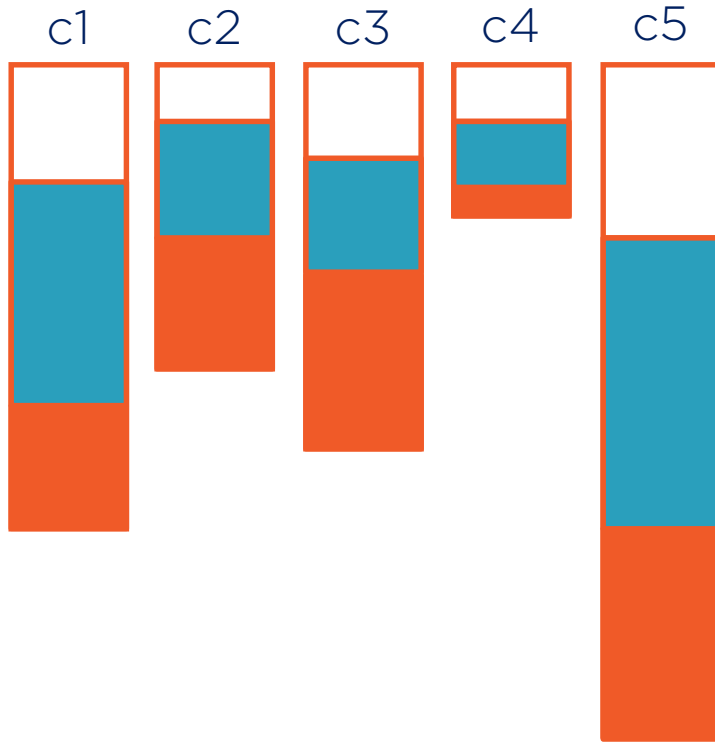
Index Structures: Row-based Indexes



Leaf (bottom) level: contains something for every table row in indexed order

Non-leaf level(s) or B-tree: contains something, specifically representing the **FIRST** value, from every page of the level below

Index Structures: Column-based Indexes



All rows for a single column are stored together – potentially HIGHLY compressible!

Data is segmented into groups of 1 million rows for better batch processing

SQL Server can do “segment elimination” (similar to partition elimination) and further reduce number of segments to process!

Parallelization through batch mode

Choice Depends on Workload: OLTP

Priority is toward
modifications

Many point queries
(highly-selective
and return very few
rows)

Row-based
clustered index
often best

Row-based
nonclustered
indexes are very
important

Might use
nonclustered
columnstore
indexes



Choice Depends on Workload: DSS/RDW

Priority is toward
large-scale
aggregates

High percentage of
or entire dataset is
evaluated often

Clustered
columnstore for
aggregates and
whole-table queries

Secondary indexes
still important

Point queries
benefit from row-
based nonclustered
indexes



Choice Depends on Workload: Hybrid

OLTP might be
priority, with some
point query activity

Some range-based
queries because
“management” wants
real-time analysis

Needs mix of
indexes
across tables

Possibly nonclustered
columnstore indexes if
data is partitioned

Best in SQL Server
2016+



DSS/RDW Strategies by Version

Prior to SQL Server 2012

- Traditional row-based clustered and nonclustered indexes

SQL Server 2012+

- Consider adding a read-only, nonclustered, columnstore index for partitioned objects leveraging partition switching as additional data is added

SQL Server 2014+

- Use SQL Server 2012+ strategy above
- Or consider the new updateable clustered columnstore index



RDW/Columnstore Indexes by Version

- **SQL Server 2008**

- Lowest version to consider for large table/data analysis with performance and scalability features
 - Added data compression (row and page compression)
 - Added filtered indexes/filtered statistics
 - Fixed fast-switching for partition-aligned, indexed views

- **SQL Server 2012**

- Added read-only, nonclustered columnstore indexes
- Some frustrating limitations but still amazing performance when possible and/or workarounds used (details at <http://bit.ly/1eHVW00>)



RDW/Columnstore Indexes by Version

- **SQL Server 2014**

- Added updateable, clustered columnstore indexes
 - Frustrating limitation that no other indexes allowed
- Many other frustrating limitations with columnstore fixed
 - E.g., UNION ALL supports batch mode, so useable with partitioned views
- Added “incremental statistics” to help limit rebuilds as well as time to rebuild

- **SQL Server 2016**

- Added updateable nonclustered columnstore indexes
- Added row-based nonclustered indexes with clustered columnstore indexes
- Any combination of base-table structure with nonclustered indexes possible



Row-based vs. Column-based Indexes

Row-based

Support data compression

Column-based

Columnar data stored together, often allows much higher level of compression



Row-based vs. Column-based Indexes

Row-based

Support data compression

Can support point queries (seeks)

Column-based

Columnar data stored together, often allows much higher level of compression

Supports large-scale aggregations



Row-based vs. Column-based Indexes

Row-based

Support data compression

Can support point queries (seeks)

Wide variety of supported scans

Full and/or partial table scans

Nonclustered covering scans

Nonclustered covering seeks with partial scans

Column-based

Columnar data stored together, often allows much higher level of compression

Supports large-scale aggregations

Support partial scans with segment elimination, and combine with partitioning for further elimination



Row-based vs. Column-based Indexes

Row-based Problems

More tuning work for analysis: must create appropriate indexes per query and then consolidate

Stores multiple columns of data together (not as easily compressed)

Column-based Problems

Minimum set for reads is a row group (no seeks possible)

Limitations of features for batch mode by version (fixes in 2014 and 2016)

Limitations with other features (fewer and fewer by SQL Server version)



Indexed Views vs. Columnstore Indexes

Indexed Views

Limited uses in non-Enterprise Editions

Columnstore Indexes

Enterprise Edition only before 2016 SP1



Indexed Views vs. Columnstore Indexes

Indexed Views

Limited uses in non-Enterprise Editions

Requires certain session settings set on

Columnstore Indexes

Enterprise Edition only before 2016 SP1

No session setting requirements



Indexed Views vs. Columnstore Indexes

Indexed Views

Limited uses in non-Enterprise Editions

Requires certain session settings set on

Must be analyzed / created “per query”

Columnstore Indexes

Enterprise Edition only before 2016 SP1

No session setting requirements

Only one can be created per table



Indexed Views vs. Columnstore Indexes

Indexed Views

Limited uses in non-Enterprise Editions

Requires certain session settings set on

Must be analyzed / created “per query”

More complicated to create

Columnstore Indexes

Enterprise Edition only before 2016 SP1

No session setting requirements

Only one can be created per table

Very easy to create



Indexed Views vs. Columnstore Indexes

Indexed Views

Limited uses in non-Enterprise Editions

Requires certain session settings set on

Must be analyzed / created “per query”

More complicated to create

More storage required

Columnstore Indexes

Enterprise Edition only before 2016 SP1

No session setting requirements

Only one can be created per table

Very easy to create

A LOT LESS storage required



Indexed Views vs. Columnstore Indexes

Indexed Views

Limited uses in non-Enterprise Editions

Requires certain session settings set on

Must be analyzed / created “per query”

More complicated to create

More storage required

**More administrative overhead /
maintenance**

Columnstore Indexes

Enterprise Edition only before 2016 SP1

No session setting requirements

Only one can be created per table

Very easy to create

A LOT LESS storage required

**Less administrative overhead /
maintenance**



Indexed Views vs. Columnstore Indexes

Indexed Views

Limited uses in non-Enterprise Editions

Requires certain session settings set on

Must be analyzed / created “per query”

More complicated to create

More storage required

More administrative overhead /
maintenance

More costly to maintain during inserts /
updates

Columnstore Indexes

Enterprise Edition only before 2016 SP1

No session setting requirements

Only one can be created per table

Very easy to create

A LOT LESS storage required

Less administrative overhead /
maintenance

Might not be able to do inserts /
updates, depending on version



Summary



Understanding ALL index types is really best as you move forward

This course will focus on row-based indexes

For columnstore indexes:

- Pluralsight: *SQL Server 2012: Nonclustered Columnstore Indexes*
- Blog series:
<http://www.nikoport.com/columnstore/>

What We Covered



Row-based index structures

Column-based index structures

Understanding your workload

Row-based vs. column-based indexes

Columnstore indexes by version

Indexed views vs. columnstore indexes

