

# Clustered Index Internals

---



**Kimberly L. Tripp**

OWNER/PRESIDENT - SQLSKILLS.COM

@kimberlyltrippp

[www.sqlskills.com/blogs/kimberly](http://www.sqlskills.com/blogs/kimberly)



# Module Overview




## Index levels

### Case study

- Building the leaf level
- Building the tree structure



# Physical Index Levels



**Leaf level:** contains something for every row of the table in indexed order

**Non-leaf level(s) :** contains something, specifically representing FIRST value, from every page of level below (intermediate levels are not a certainty)

# Employee Table Case Study

## Clustering the Employee table on EmployeeID (identity)

- Investigate the data
- Physically order data
- Add the tree structure
- Complete the math
- Complete the B-tree

## Nonclustered unique constraint for SSN

- In the next module...



# Investigating the Data

**Average row size  
= 400 bytes**

```
CREATE TABLE [Employee]
(
  [EmployeeID]      int          NOT NULL identity,
  [LastName]        nchar(30)    NOT NULL,
  [FirstName]       nchar(29)    NOT NULL,
  [MiddleInitial]   nchar(1)     NULL,
  [SSN]             char(11)     NOT NULL,
  ...other columns...)
```

**Header 96 bytes**

8,096 bytes  
per page  
for data



$$\frac{8,096 \text{ bytes / page}}{400 \text{ bytes / row}} = 20 \text{ rows/page}$$

$$80,000 \text{ current employees} = 80,000 \text{ rows} \quad \frac{80,000 \text{ employees}}{20 \text{ rows / page}} = 4,000 \text{ pages}$$


# Clustered Employee Table

## Step 1: Physically order data

Review the index level definitions...  
Does this seem to match one of the definitions?

**Yes!** When a table is clustered the data becomes the leaf level of the clustered index!

4,000 pages of Employees in clustering key (EmployeeID) order



# Clustered Employee Table

## Step 2: Add the tree structure

Starting from the leaf level and going up to a root of 1 page

**B-tree entry = index key value + pointer + row overhead**

Pointer = page pointer of 6 bytes = 2 for fileID + 4 for pageID

Row overhead varies based on many factors (min of 1 byte in the row)

**Non-leaf level entry for clustered index on EmployeeID = 11**

4 bytes for EmployeeID (int) + 6 bytes for page pointer + 1 byte for row overhead

$$\frac{8,096 \text{ bytes / page}}{11 \text{ bytes/entry} + 2 \text{ bytes in slot array}} = 622 \text{ index entries per non-leaf level page}$$

How many entries to store? **4,000**

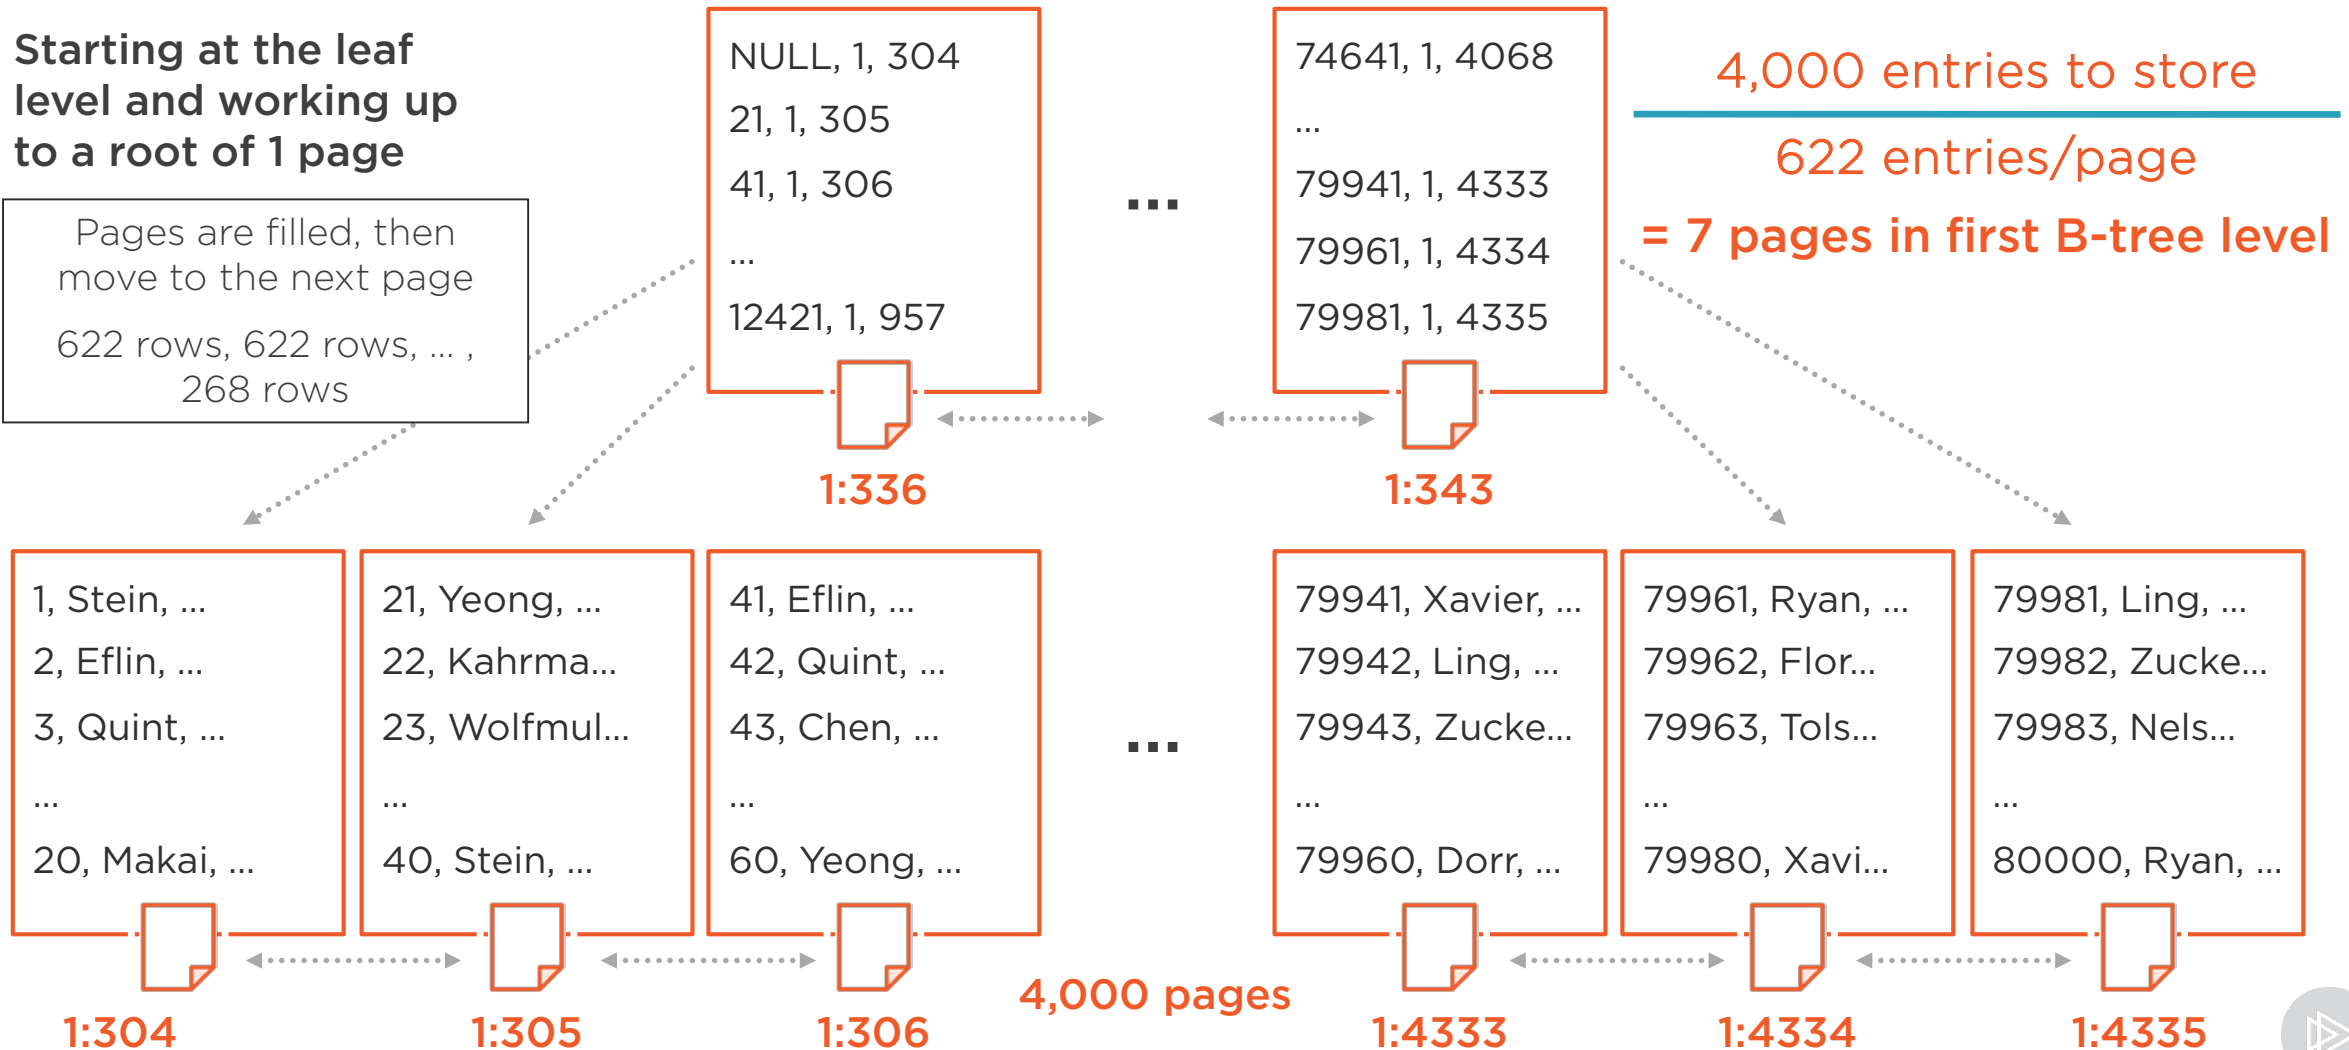


# Clustered Employee Table

## Step 2: Add the tree structure

Starting at the leaf level and working up to a root of 1 page

Pages are filled, then move to the next page  
622 rows, 622 rows, ... ,  
268 rows





# Clustered Employee Table

## Step 2: Add the tree structure

Continuing up to  
a root of 1 page

Root  
= 1 page

NULL, 1, 336  
12441, 1, 337  
24881, 1, 339  
...  
74641, 1, 343

1:338

Intermediate level  
= 7 pages

1:336

1:343

B-tree total  
overhead in terms  
of disk space  
= 8 pages  
or < 1%

Leaf level  
= 4,000 pages

1:304

1:305

1:306

1:4333

1:4334

1:4335



# Demo



## Clustered index internals



# What We Covered



## Index levels

### Case study

- Building the leaf level
- Building the tree structure

