

Covering Queries



Kimberly L. Tripp

OWNER/PRESIDENT - SQLSKILLS.COM

@kimberlyltrippp

www.sqlskills.com/blogs/kimberly



Module Overview



Nonclustered index seeks and scans

Selectivity determines access pattern

Understanding covering

How is covering possible

Improving low-selectivity queries

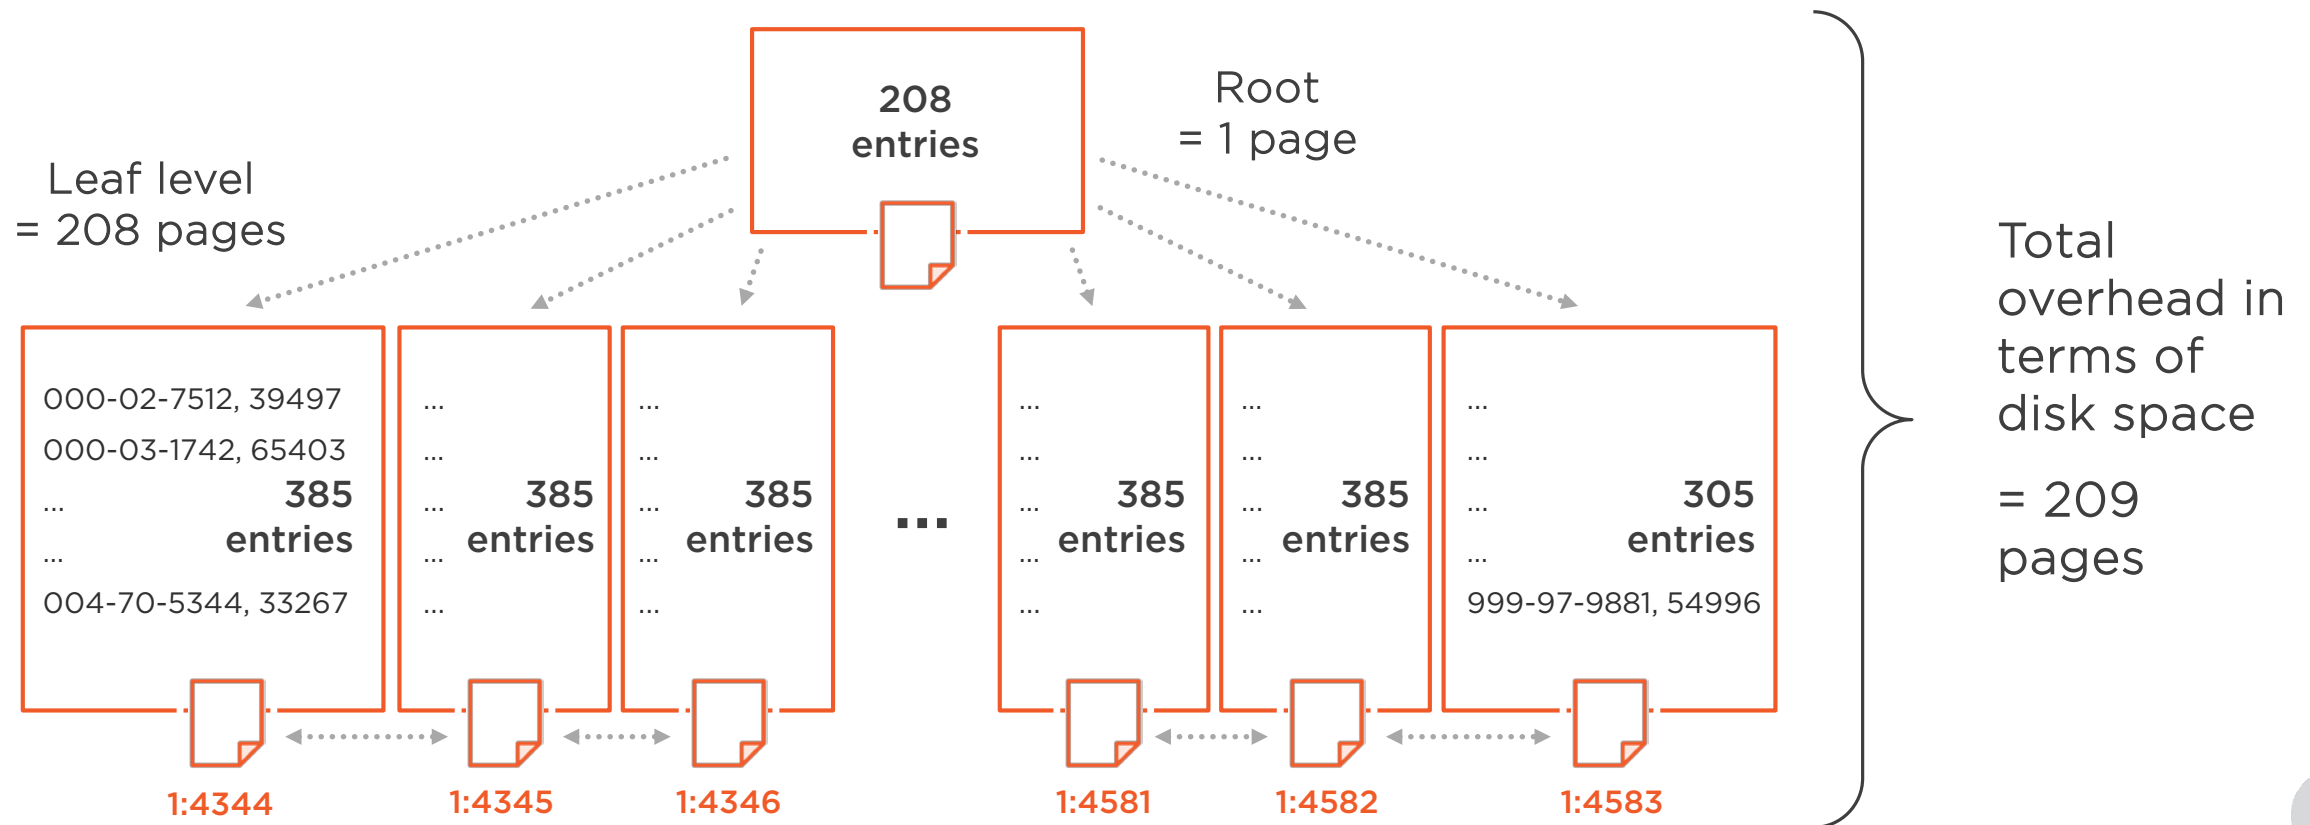
Covering cleverly, correctly, and
concisely!



Nonclustered Index: Unique Key SSN

Think back to our nonclustered index structure

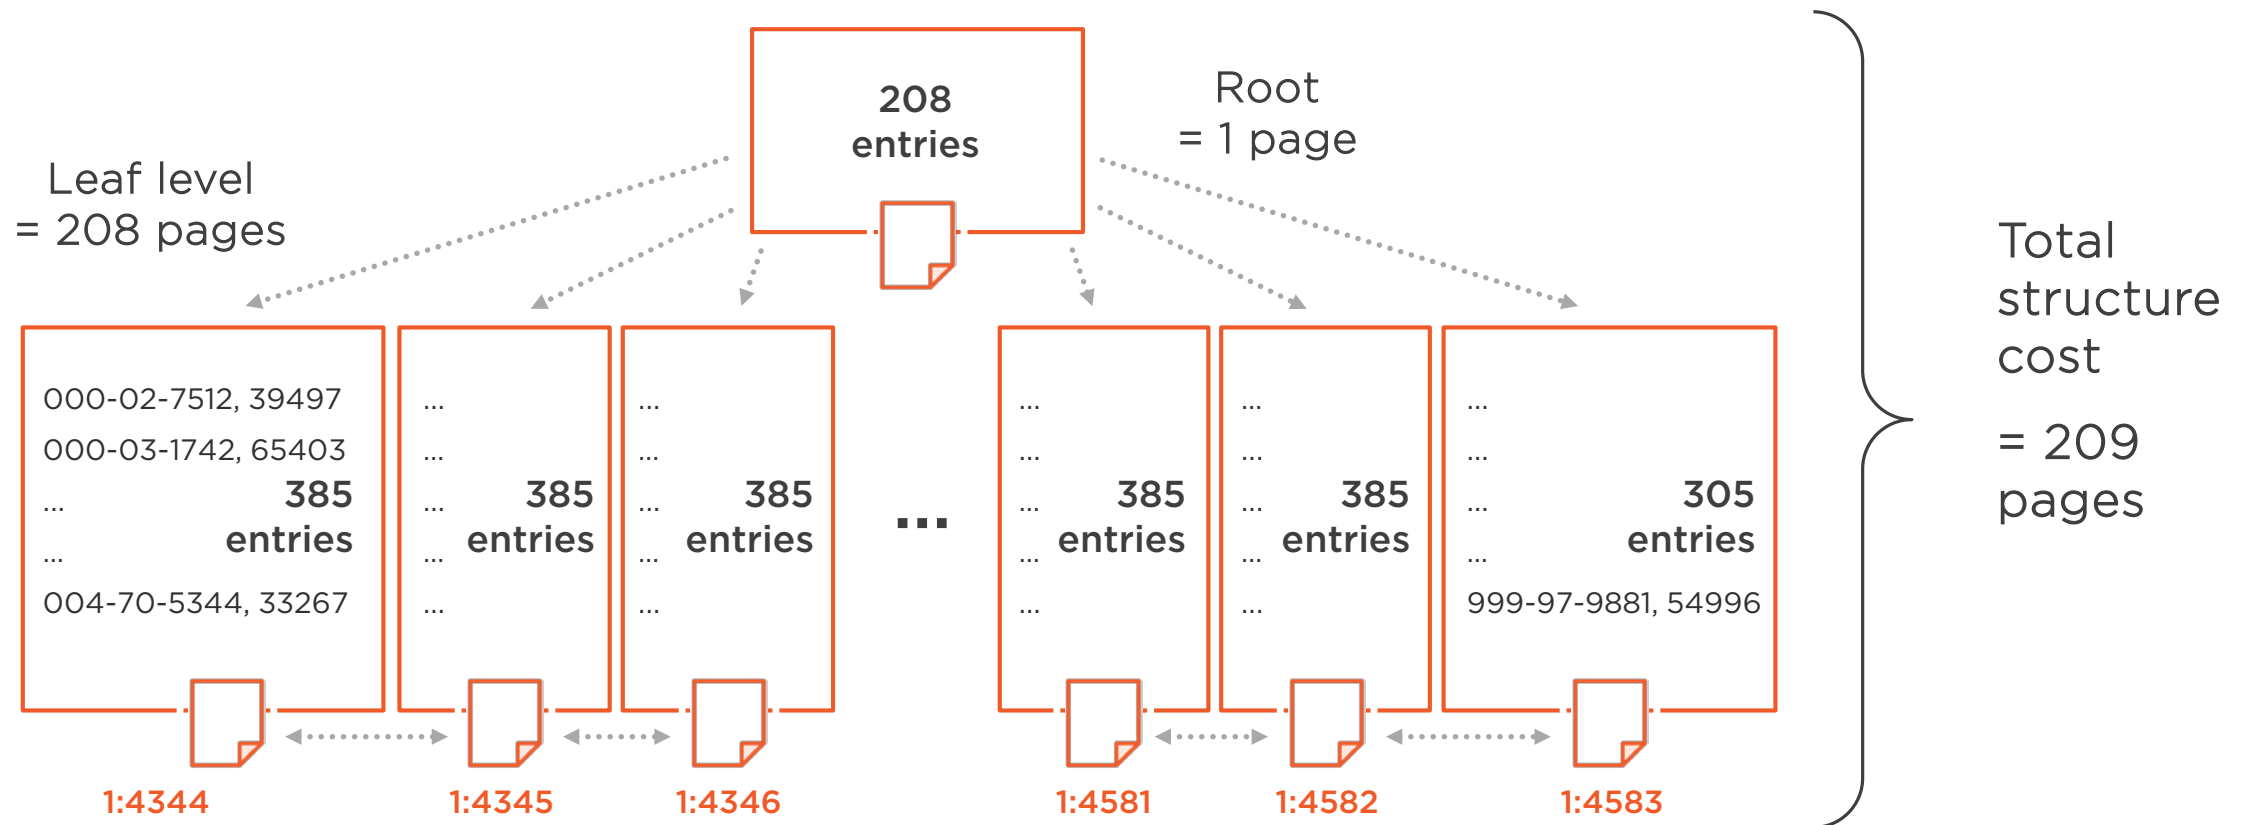
Leaf level contains nonclustered key column(s), in indexed order



What if You Didn't Know?

Could this structure be anything else?

What about table with only two columns (EmployeeID and SSN)
that's clustered on SSN?

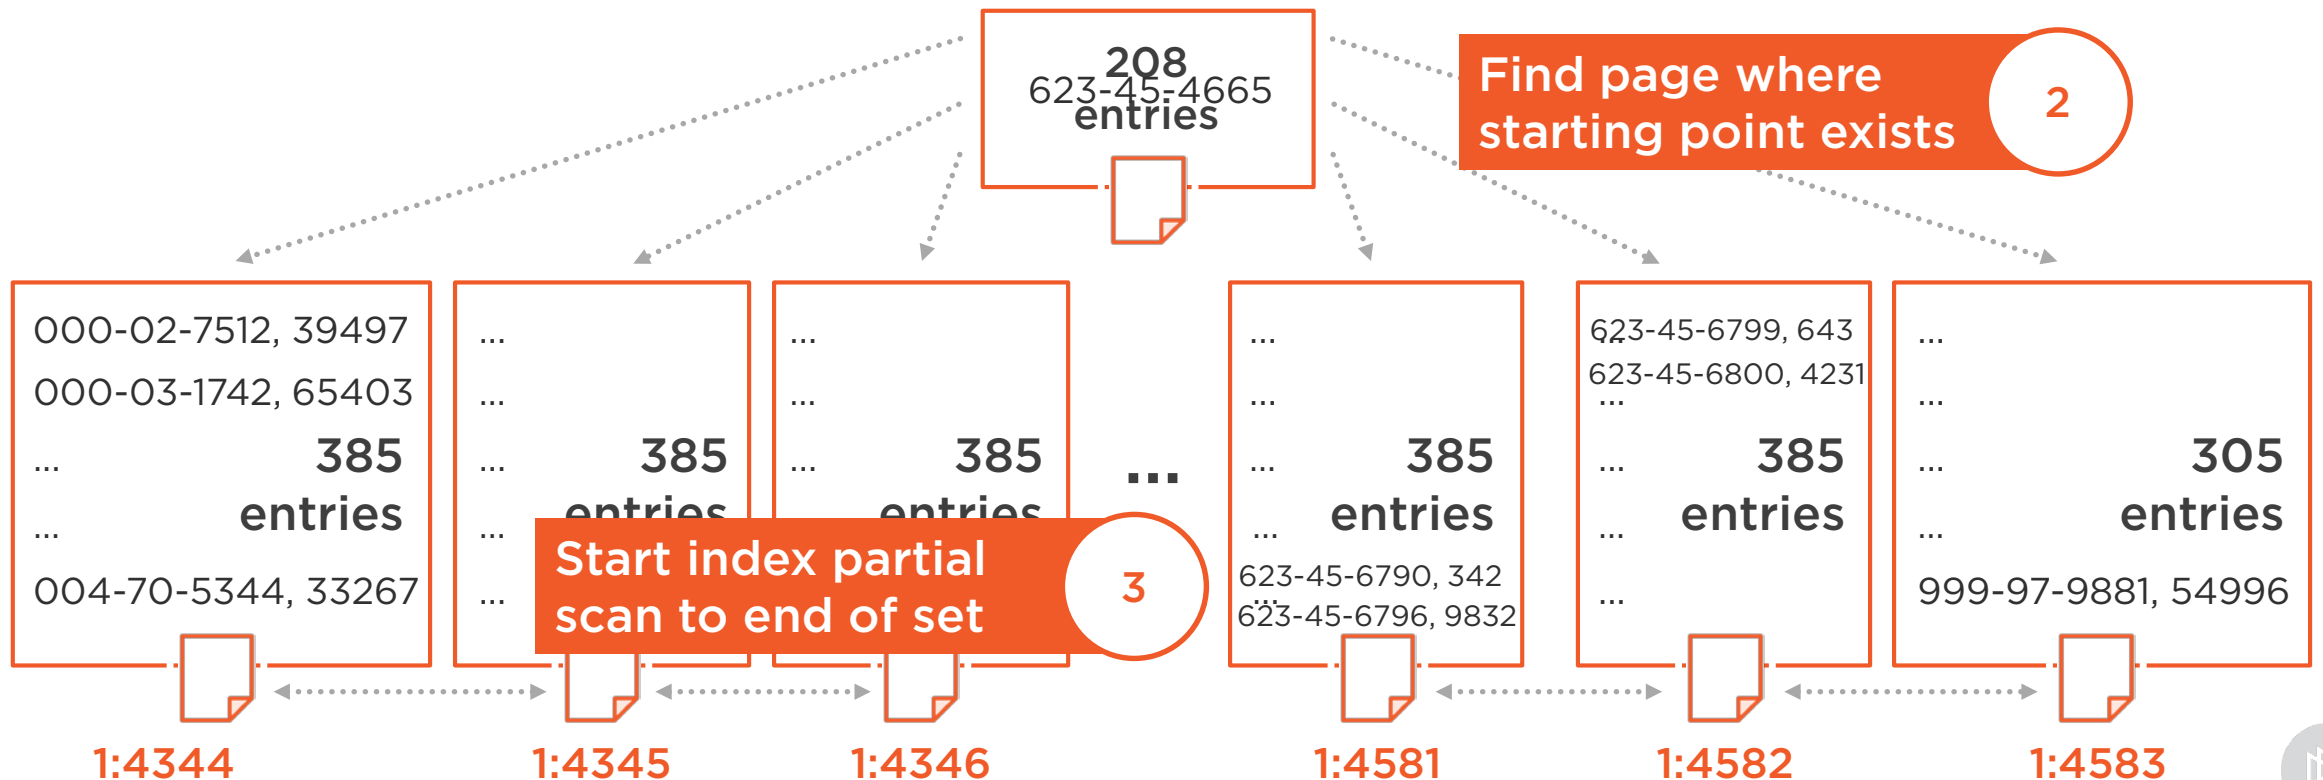


Fairly Obvious Index Access

```
SELECT [e].[EmployeeID], [e].[SSN]
FROM [dbo].[Employee] AS [e]
WHERE [e].[SSN] BETWEEN '623-45-6789'
      AND '623-45-6800';
```

Start at root page 1

Find page where starting point exists 2



Demo



Nonclustered, covering, seek



Less Obvious Index Access

```
SELECT [e].[EmployeeID], [e].[SSN]  
FROM [dbo].[Employee] AS [e]  
WHERE [e].[EmployeeID] < 10000;
```

Clustered index on EmployeeID = ~500 reads

- Seekable with partial scan
- Table has 80,000 rows at 20 rows per page, so 4,000 pages
- 10,000 is $\frac{1}{8}$ of 80,000, so this SEEKABLE query will cost $\frac{1}{8}$ of 4,000 pages

Nonclustered index on SSN, EmployeeID = ~208 reads

- Not seekable, must scan
- Leaf level has 80,000 rows at 385 rows per page, so 208 pages
- If index is not seekable, then SQL Server must scan all 208 pages



Demo



Nonclustered, covering, scan



The Best Index Varies...

What if number of rows **WHERE EmployeeID < 10000** was not 9,999?

Would optimizer make a different choice between the two indexes?

- When 9,999 rows:
 - Seekable clustered index = ~500 reads
 - Nonclustered covering index scan = ~208 reads
- If only 1,000 rows, for example:
 - Seekable clustered index = ~50 reads
 - Nonclustered covering index scan = ~208 reads

Because the covering index isn't seekable, the best index varies!

If query is critical, consider covering with a nonclustered index that is seekable:

```
CREATE INDEX [NCCoveringSeekableIX]  
ON [dbo].[Employee] ([EmployeeID], [SSN])
```



Demo



Which index is best varies



What Is Covering?

Only applies to nonclustered indexes

- Clustered “covers” all requests but it’s what we’re trying to avoid!

Using JUST a nonclustered index to access the data a query requests

All columns requested in query are somewhere in index regardless of:

- Where they are in the query
- Where they are in the index

However, column order does matter...

- If an index is seekable, might be able to drastically reduce reads
- If an index is not seekable and can only be scanned, still might have significant savings, depending on how much narrower index is than base table...



How Is Covering Possible?

Unfiltered NC indexes contain a row for every row in base table

Each NC is a “mini clustered table” of just the data you need!

This structure looks and acts just like a table

We can seek into a NC index; we can scan a NC index




We can AVOID bookmark lookups!

There are all sorts of ways we can leverage this!



Improve Critical Queries with Low Selectivity

```
SELECT [e].[LastName], [e].[FirstName], [e].[Phone]
FROM [dbo].[Employee] AS [e]
WHERE [e].[LastName] LIKE '[S-Z]%' ;
-- 13,632 in S-Z range (NOT selective)
```



Options to Access Data

Table scan

**Nonclustered on LastName
with bookmark lookups for
every row**

**Nonclustered index:
LastName, FirstName,
PhoneNo**

**Nonclustered index:
FirstName, LastName,
PhoneNo**



You CANNOT Cover Everything!

Theoretically you can cover anything and everything

- *Just because you can, doesn't mean you should!*

Over-indexing can be worse than under-indexing... as unused indexes are nothing but costly overhead

- *Do not just automatically put “an index on every column”!!*

Too many indexes cost you:

- During data modifications
- During regular index maintenance
- Wasted space in memory
- Wasted space on disk (including in backups)



Covering: Cleverly, Correctly, and Concisely!

Scalability != “add more indexes”

Get a good feel for your workload

Know your data and workload, plus how SQL Server works

Test a wide variety of different values and scenarios

Using covering sparingly

Use index consolidation techniques



Methods for Covering

**Nonclustered
indexes**

Using INCLUDE

**Using filtered
indexes**

**Using indexed
views (more for
relational DW)**

**Using columnstore
indexes (more for
relational DW)**



What We Covered



Nonclustered index seeks and scans

Selectivity determines access pattern

Understanding covering

How is covering possible

Improving low-selectivity queries

Covering cleverly, correctly, and concisely!

