

SQL Server: Introduction to Extended Events

Module 9: Basic Troubleshooting Scenarios

Jonathan M. Kehayias
Jonathan@SQLskills.com



Introduction

- Like SQL Trace, Extended Events should be used in situations where they provide the best method of identifying problems
- Some basic troubleshooting scenarios where Extended Events can be used include:
 - Identifying deadlock issues
 - Identifying blocking issues
 - Identifying recompilation issues
 - Identifying when errors occur
 - Tracking session-level waits

Identifying Deadlock Issues

- The `xml_deadlock_report` event fires when the Lock Monitor in SQL Server identifies a deadlock and raises error 1205 (when killing the deadlock victim process/processes)
- XML report that contains all of the information needed to troubleshoot deadlocks
- **New XML format in Extended Events**
 - Supports multi-victim deadlock analysis
 - Incompatible with graphical display of deadlock graph in SSMS
 - Reduces redundant information that existed in previous XML format
- **The RTM releases of SQL Server 2008 and 2008R2 contain a bug which causes the new XML format to be incorrectly formed**
 - SQL Server 2008 SP1+CU6 or higher or SQL Server 2008R2 RTM+CU1 or higher fix this bug (<http://support.microsoft.com/kb/978629>)
- **Collected by default on all instances of SQL Server 2008 onwards in the `system_health` session**

Identifying Blocking Issues

- The `blocked_process_report` event fires based on the value configured for the 'blocked process threshold' `sp_configure` option in the SQL Server
- XML report that contains information about the blocking and blocked processes in a blocking scenario for further debugging to identify and prevent the problem
- Setting the 'blocked process threshold' too low can result in excessive event generation
 - For example, if the threshold is set at 10 seconds and a blocking scenario lasts for 38 seconds, three `blocked_process_report` events will be generated (one every 10 seconds)
 - In the same example, if there are multiple blocked sessions in a blocking chain, each blocked session will generate a `blocked_process_report` event every 10 seconds

Identifying Recompilation Issues

- **The `sql_statement_starting` and `sp_statement_starting` events contain a 'state' column that specifies whether the statement was recompiled during execution**
 - The state column is mapped to the `statement_starting_state` map and provides three values: Normal, Recompiled, and Execution Plan Flush
 - Recompilation causes the event to fire twice: once for `state=Recompiled` and once for `state=Normal`
- **The `sql_statement_recompile` event fires for any statement-level recompilation in the system**
 - Ad hoc batches, stored procedures, and triggers are included
 - The `recompile_cause` column is mapped to the `statement_recompile_cause` map and provides the reason the recompile occurred

Identifying When Errors Occur

- **The error_reported event fires when errors occur during execution, even if the error is handled by Transact-SQL code**
 - The is_intercepted column can be used to determine if the error was handled by a TRY/CATCH block in the Transact-SQL code
- **Using filters on the error_number, severity, and/or state can make it easier to identify the root cause of specific problems**
 - 220 - Arithmetic overflow error for data type...
 - 8152 - String or binary data would be truncated...
 - 2627 - Cannot insert duplicate key in object...
- **Actions like the tsql_stack, query_hash, and sql_text can make identifying the point in code where the error occurred possible**
 - Prevents collecting sql_statement_starting and/or sp_statement_starting events to correlate where the error is occurring

Tracking Session Wait Statistics

- Understanding the causes of waits inside SQL Server can help identify performance bottlenecks and potential future problems
- The `wait_info` and `wait_info_external` events fire whenever a task has to wait during its execution
- Predicates on the `session_id` global field can allow tracking waits for a specific session in the server, or can be used to sample all sessions on the server
- An example of how to track session level waits is in Module 3 of the [SQL Server: Performance Troubleshooting Using Wait Statistics](#) course

Viewing Historical System Health

- In SQL Server 2012, `sp_server_diagnostics` executes in five minute intervals to check the current health of the SQL Server instance
- It provides four sections of health information:
 - System – spinlock backoffs, CPU usage, memory dumps, non-yielding tasks
 - Resource – process and memory manager memory usage
 - Query Processor – max workers, current number of workers, idle workers, tasks completed in interval, top 10 non-preemptive and preemptive wait types by count
 - I/O subsystem – latch timeouts, long interval I/O, pending requests
- The output of `sp_server_diagnostics` is collected by the `system_health` session and written to a file in the file system
 - Data can be trended over longer periods of time
 - Data will be available if a crash or restart occurs for the instance

Summary

- **Basic troubleshooting can be accomplished using Extended Events**
- **In some cases the data necessary for troubleshooting problems, like deadlocks occurring, is already being collected by the system_health event session on the server**

- **Many more scenarios for using Extended Events exist than were covered in this module**
 - These are merely suggested scenarios where basic troubleshooting can be performed using Extended Events
 - More advanced scenarios will be covered by the forthcoming Advanced Extended Events course