

# Diagnosing Activity-related Issues

---



**Glenn Berry**

PRINCIPAL CONSULTANT - SQLSKILLS.COM

@GlennAlanBerry

[www.sqlskills.com/blogs/glenn](http://www.sqlskills.com/blogs/glenn)



# Module Overview



**Diagnosing activity-related issues**

**Interpreting query results**

**Alleviating activity-related issues**



# Lock Waits

Tables and  
indexes that  
have lock waits

Very useful if  
there are very high  
average task counts

Shows row  
lock waits

Shows page  
lock waits

Cumulative waits  
since last restart

Index tuning  
can often reduce  
lock waits



# Lock Waits



Look for tables that have high row and/or page lock waits



Look for example of high lock waits on clustered index of a table



Adding useful non-clustered indexes can help reduce lock waits



Dropping unused indexes can help reduce lock waits



Isolation level properties have a role in concurrency/locking/blocking issues



# Demo



## Lock waits



# Scalar UDF Statistics

Scalar UDF  
metrics for  
current database

Results are  
ordered by total  
worker time

Helps find most  
CPU-intensive  
scalar UDFs



# Scalar UDF Statistics



Scalar UDFs have known performance issues in SQL Server 2017 and earlier



Consider in-lining scalar UDF code if possible



Can convert to table-valued UDF that returns one column and row



Alternative is to convert scalar UDF to a T-SQL stored procedure



# Demo



## Scalar UDF statistics





# Input Buffer

**Replacement  
for DBCC  
INPUTBUFFER**

**Shows last query  
for each SPID  
connected to  
current database**

**Returns useful  
performance  
metrics for  
each SPID**



# Input Buffer



More capable and flexible replacement for DBCC INPUTBUFFER



Use for getting a quick overview of current query workload



Can add ORDER BY clause to focus on one specific area



Helps identify resources of long running queries that are still executing



# Demo



## Input buffer



# Query Execution Counts

Most frequently  
executed queries  
for current  
database

Look for  
“Has Missing  
Index” column

Look at graphical  
execution plan



# Query Execution Counts



Helps understand baseline query workload



Frequently executed queries may be candidates for middle-tier or client-side caching



Extremely high counts may indicate application logic issues



Helps identify possible query and index tuning opportunities



# Demo



## Query execution counts



# SP Execution Counts

**Most frequently  
executed SPs for  
current database**

**Look for  
“Has Missing  
Index” column**

**Look at graphical  
execution plan**



# SP Execution Counts



Helps you understand your baseline stored procedure workload



Frequently executed SPs may be candidates for middle-tier or client-side caching



Extremely high SP counts may indicate application logic issues



Helps you identify possible SP and index tuning opportunities





# Demo



## SP execution counts



# SP Avg Elapsed Time

Cached stored  
procedures ordered  
by average  
execution time

Elapsed times are  
in microseconds

Look for large  
differences between  
min and max

Look for  
“Has Missing  
Index” column

Look at the  
graphical  
execution plan



# SP Avg Elapsed Time



Helps identify possible easy tuning opportunities



Wide variance in execution times can indicate plan stability problems



Focus your initial tuning efforts on top five results



Dramatically reducing elapsed time of a stored procedure is very beneficial!



# Demo



SP avg elapsed time



# Bad Nonclustered (NC) Indexes

Returns NC indexes that have more writes than reads

Consider dropping these indexes after more analysis

Make sure you know how long instance has been running



# Bad Nonclustered Indexes



Indexes with far more writes than reads may not be useful for workload



Azure SQL Database must update these indexes as data changes



Unused indexes increase database size and maintenance workload



Make sure you have seen your complete workload before dropping indexes



# Demo



## Bad nonclustered indexes



# Missing Indexes

Missing indexes  
for current  
database

This query is very  
useful but easy  
to misinterpret

Do careful analysis  
before adding  
new indexes





# Missing Indexes



Look at all of the columns returned by this query



Know how long database has been running as you interpret results



Pay special attention to “last\_user\_seek”, “user\_seeks”, and “avg\_total\_user\_cost” columns



Consider existing indexes and try to create fewer, wider indexes



# Demo



## Missing indexes



# Missing Index Warnings

**Finds missing  
index warnings  
in plan cache**

**Query can take  
a long time to  
return results**

**Returns object  
name and  
query plan**



# Missing Index Warnings



Can associate missing index requests with specific stored procedures



“Usecounts” column shows count of times index was requested by SP/query



Execution plan will have missing index details



Knowing which stored procedure is generating the request helps you make better at tuning decisions



# Demo



## Missing index warnings



# Overall Index Usage – Reads

Shows which  
indexes have the  
most reads

Helps you  
understand  
your workload

Index reads are  
beneficial for  
**SELECT** query  
performance



# Overall Index Usage – Reads



Indexes with high reads may benefit from data compression



Evaluate data volatility and compressibility



Tables with extremely high reads might be columnstore index candidates



Returns cumulative metrics for all row-store indexes in current database



# Demo



Overall index usage – reads





# Overall Index Usage – Writes

**Shows which  
indexes have the  
most writes**

**Helps you  
understand  
your workload**

**Index writes are  
bad for INSERT /  
UPDATE query  
performance**



# Overall Index Usage – Writes



Look for indexes with many more writes than reads



Make sure you know how long database has been running



Do not blindly drop indexes without more analysis



Returns cumulative metrics for all row-store indexes in current database



# Demo



Overall index usage – writes



# Volatile Indexes

Shows which  
indexes and  
statistics have  
most updates

Helps you  
understand your  
write workload

Helps you design  
and configure  
your storage



# Volatile Indexes



Be more cautious about creating new indexes on volatile tables



Be more cautious about using data compression on volatile tables



Consider moving highly volatile tables/indexes to separate file group



Consider using flash storage or non-parity RAID levels for volatile data



# Demo



## Volatile indexes



# Recent Resource Usage

Snapshot every 15  
seconds

Data goes back  
64 minutes

Shows average  
percentages

CPU and memory  
usage

Data and log file  
activity

Validates your  
service tier



# Recent Resource Usage



Snapshot of average values every 15 seconds for past 64 minutes



Shows average CPU percent and average memory usage percent



Shows average data file IO percent and average log write percent



These percentages are against the maximum limit for your service tier



If you are regularly near 100%, you need to tune or go to a higher service tier





# Demo



## Lock waits



# What We Covered



**Diagnosing activity-related issues**

**Interpreting query results**

**Alleviating activity-related issues**

