

# Understanding Query Optimization in SQL Server

## Introducing Query Optimization in SQL Server



**Nikola Ilic**

Data Mozart

@DataMozart | [www.data-mozart.com](http://www.data-mozart.com)

# Overview



## “Good enough” performance

- Performance tuning process
- Most common bottlenecks

## SQL Server Query Optimizer

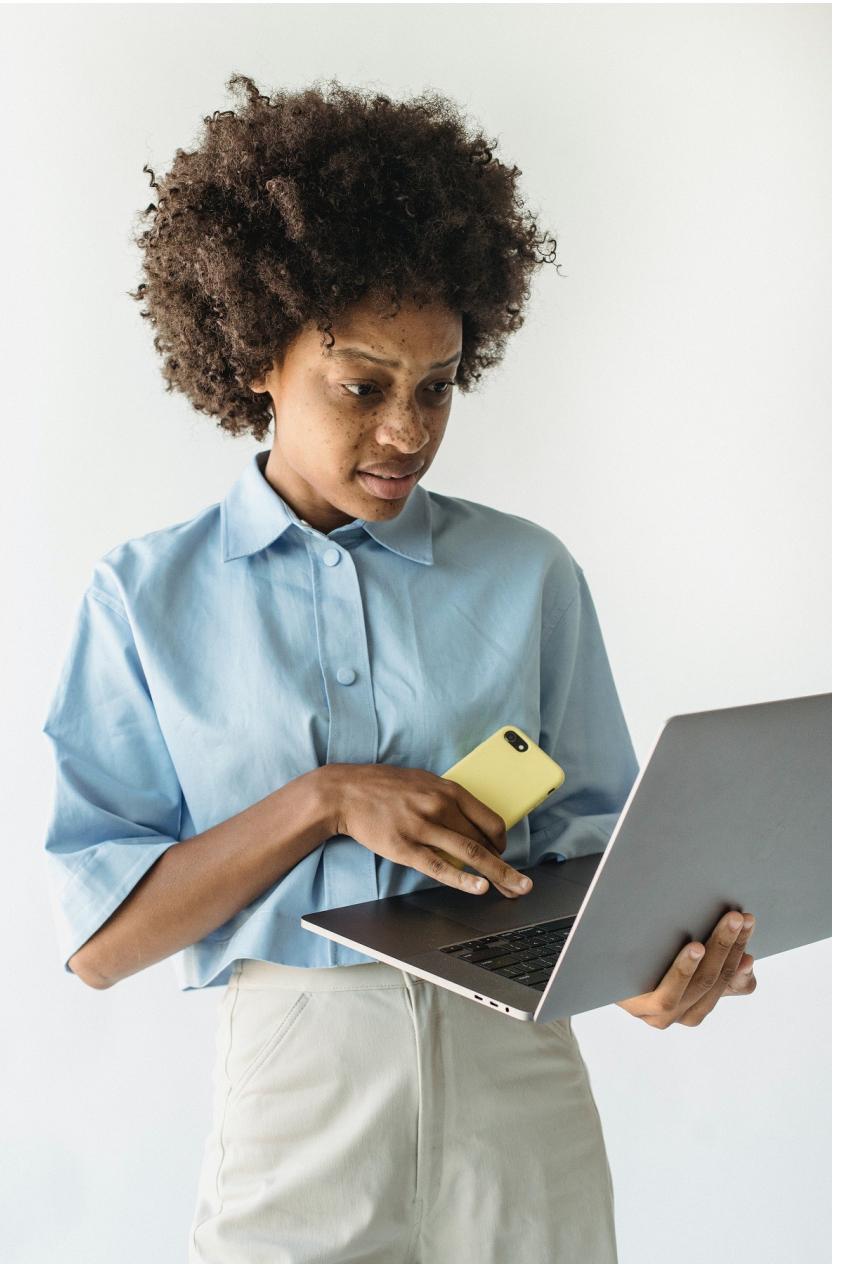
- T-SQL statement evaluation order

## Capture query performance metrics



# Is 30 seconds a lot or not?





**Buying tickets for your favorite music band**



**Waiting for the bus**

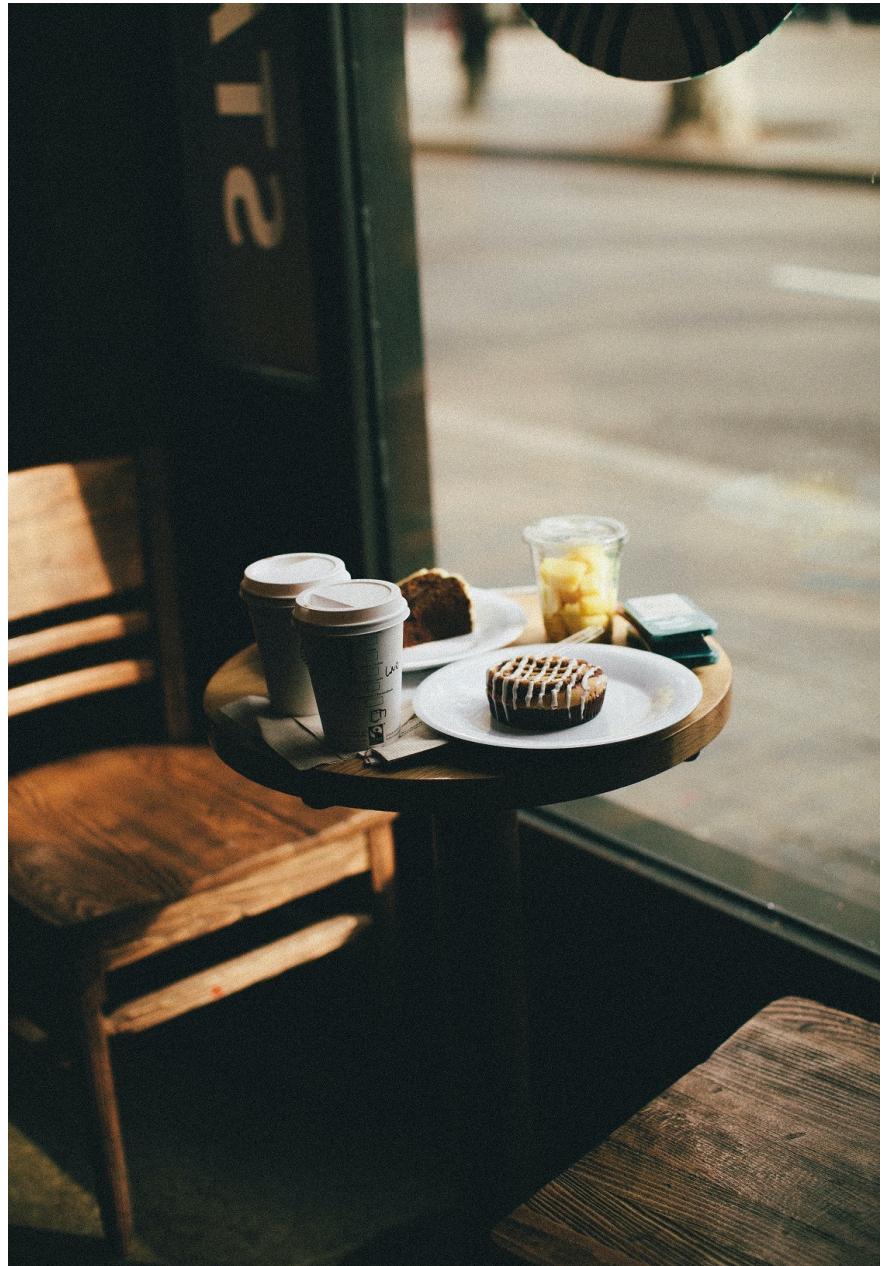


# Is 10 seconds a lot or not?





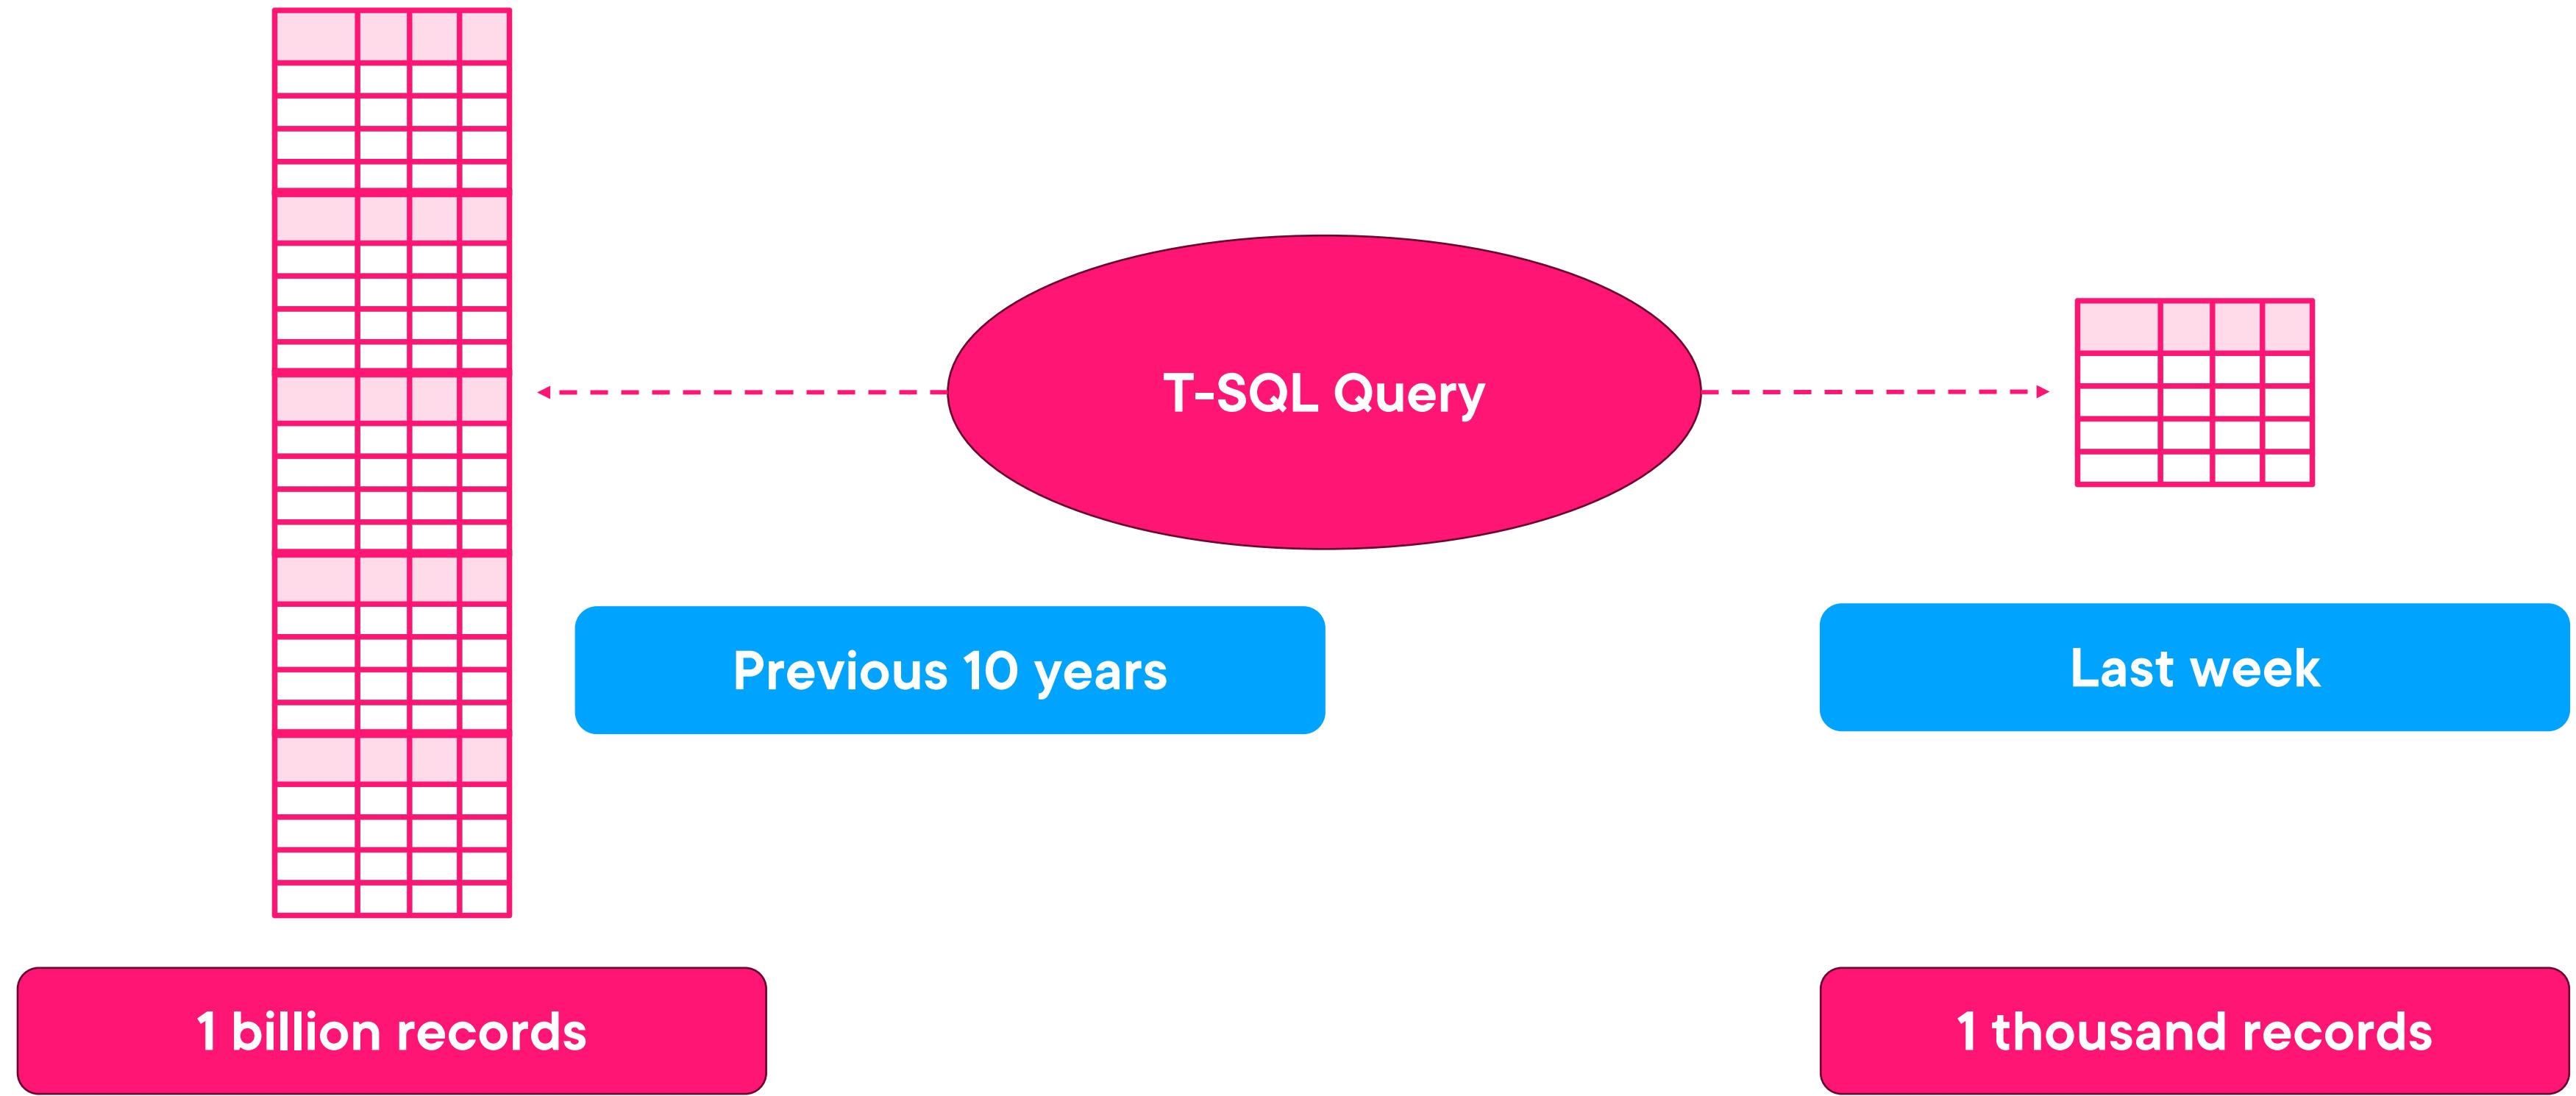
**Waiting for breaking news to render**



**Getting the coffee in Starbucks**



# Setting the Right Expectations

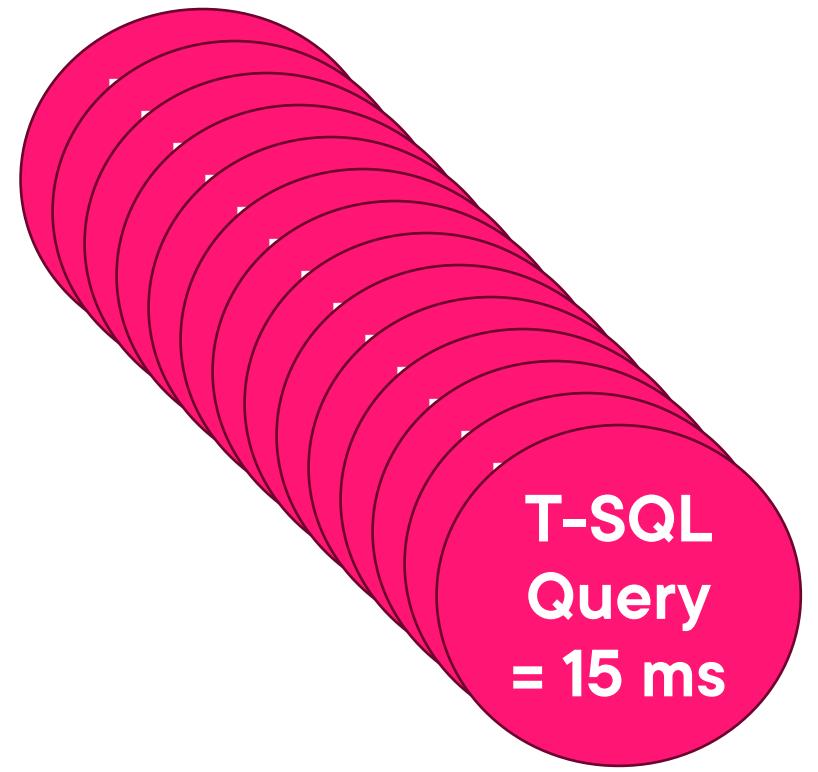


# When should you start tuning queries?

*When they are not performing  
GOOD ENOUGH...*



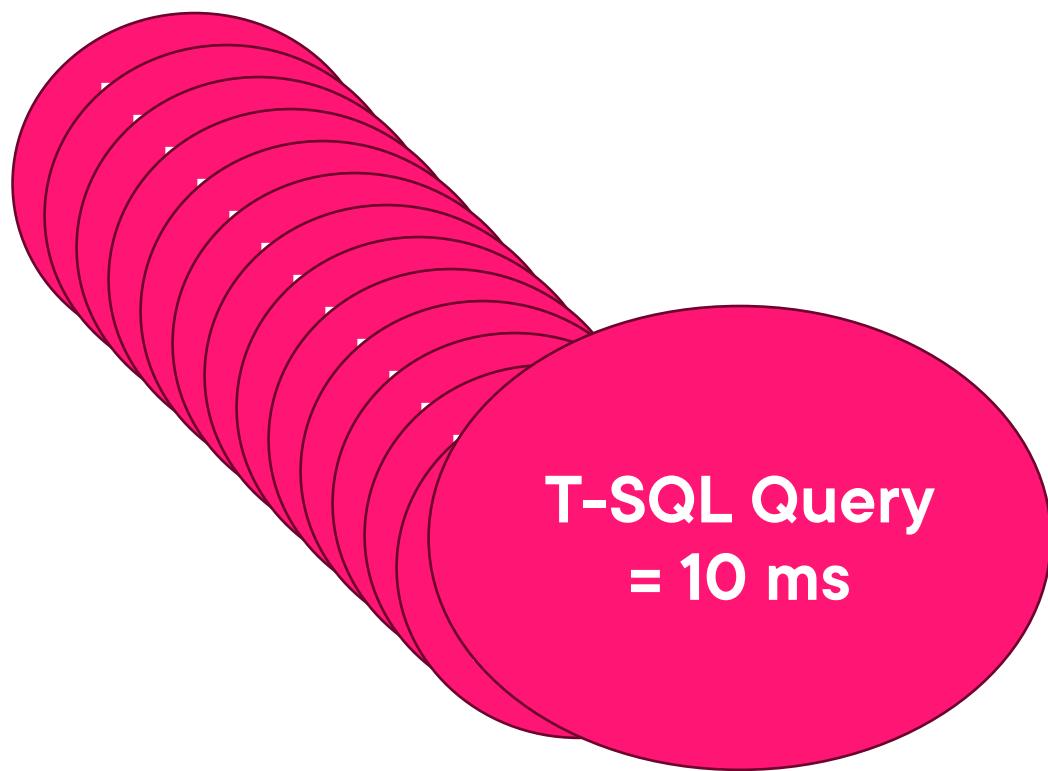
# What Is “Good Enough”?



1000x/minute



# What Is “Good Enough”?



T-SQL Query  
= 10 ms

1000x/minute

T-SQL Query  
= 50 ms

5x/hour

Is it worth reducing  
to 40ms?!



# Validation Points in Query Tuning Process

1

Performance of the app or reports that rely on database queries

2

Checking the query plan cache

3

Running queries on non-production system



# Validation Points in Query Tuning Process

1

Performance of the app or reports that rely on database queries

2

Checking the query plan cache

3

Running queries on non-production system

Time

CPU

Disk I/O



# The Most Common Performance Bottlenecks

Inadequate indexing strategy

Missing/Inaccurate Statistics

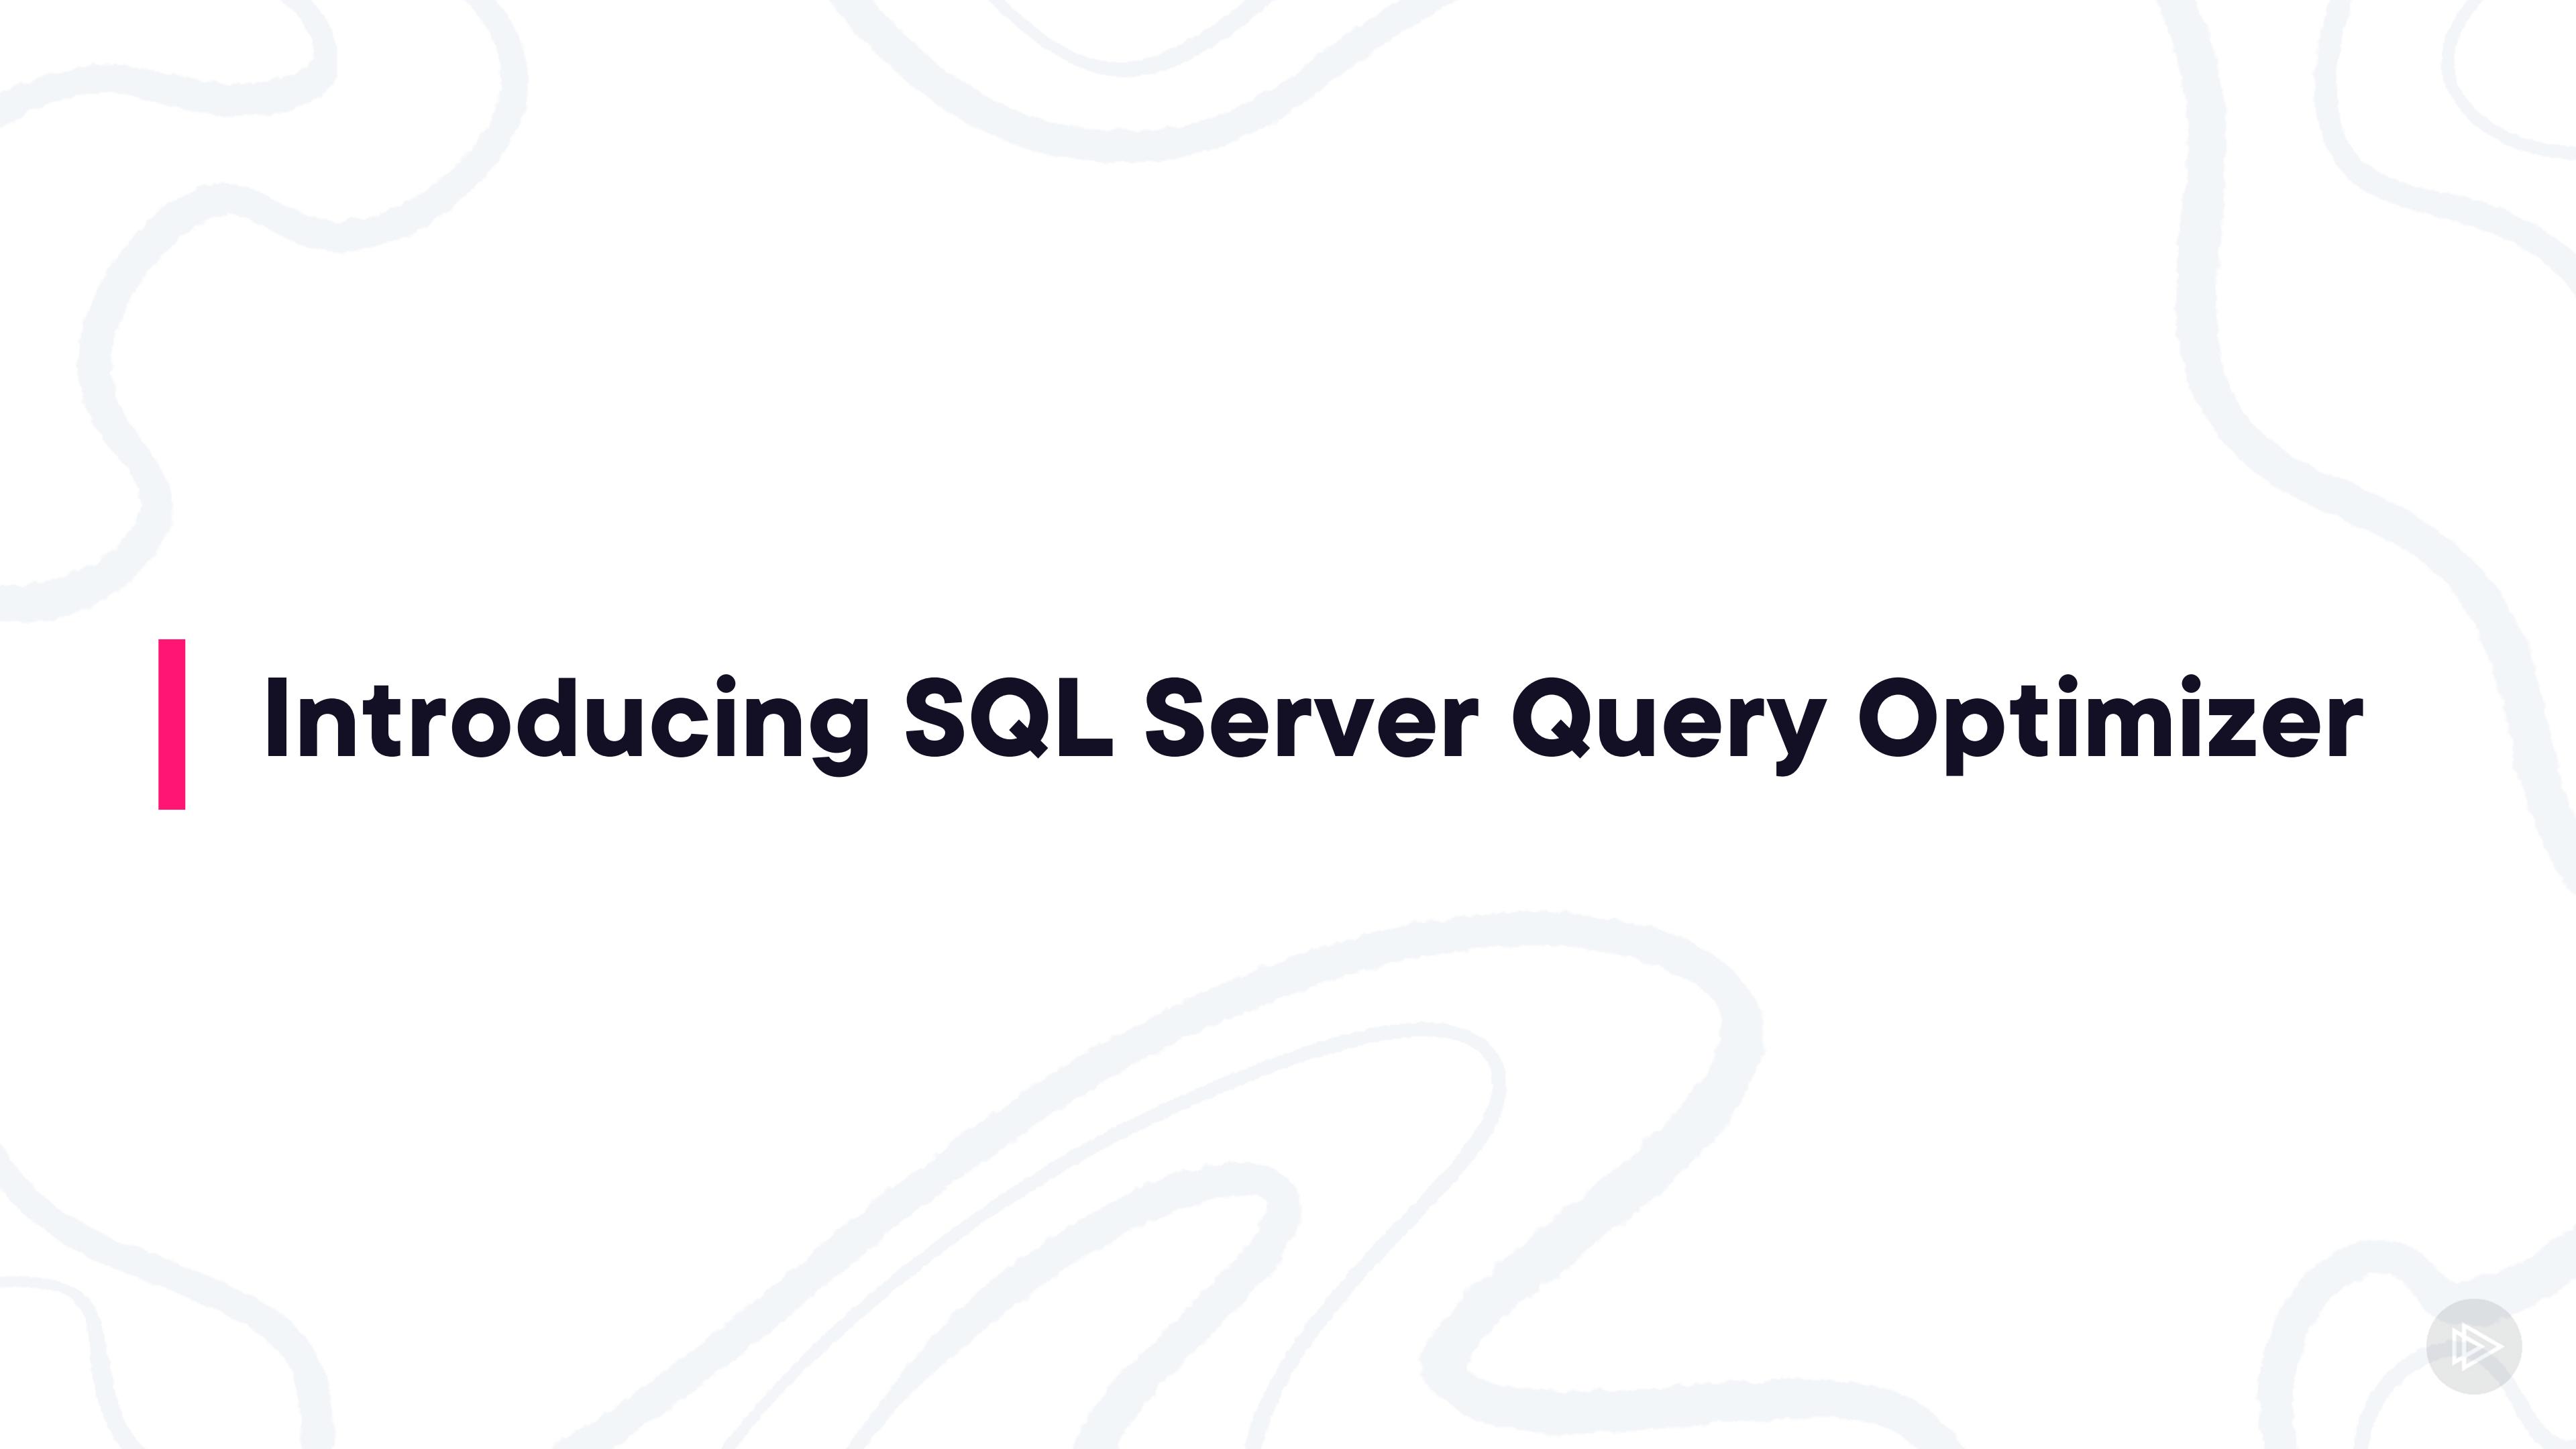
Sub-optimal T-SQL code

Sub-optimal execution plans

Sub-optimal execution plan reuse

Frequent query recompilation



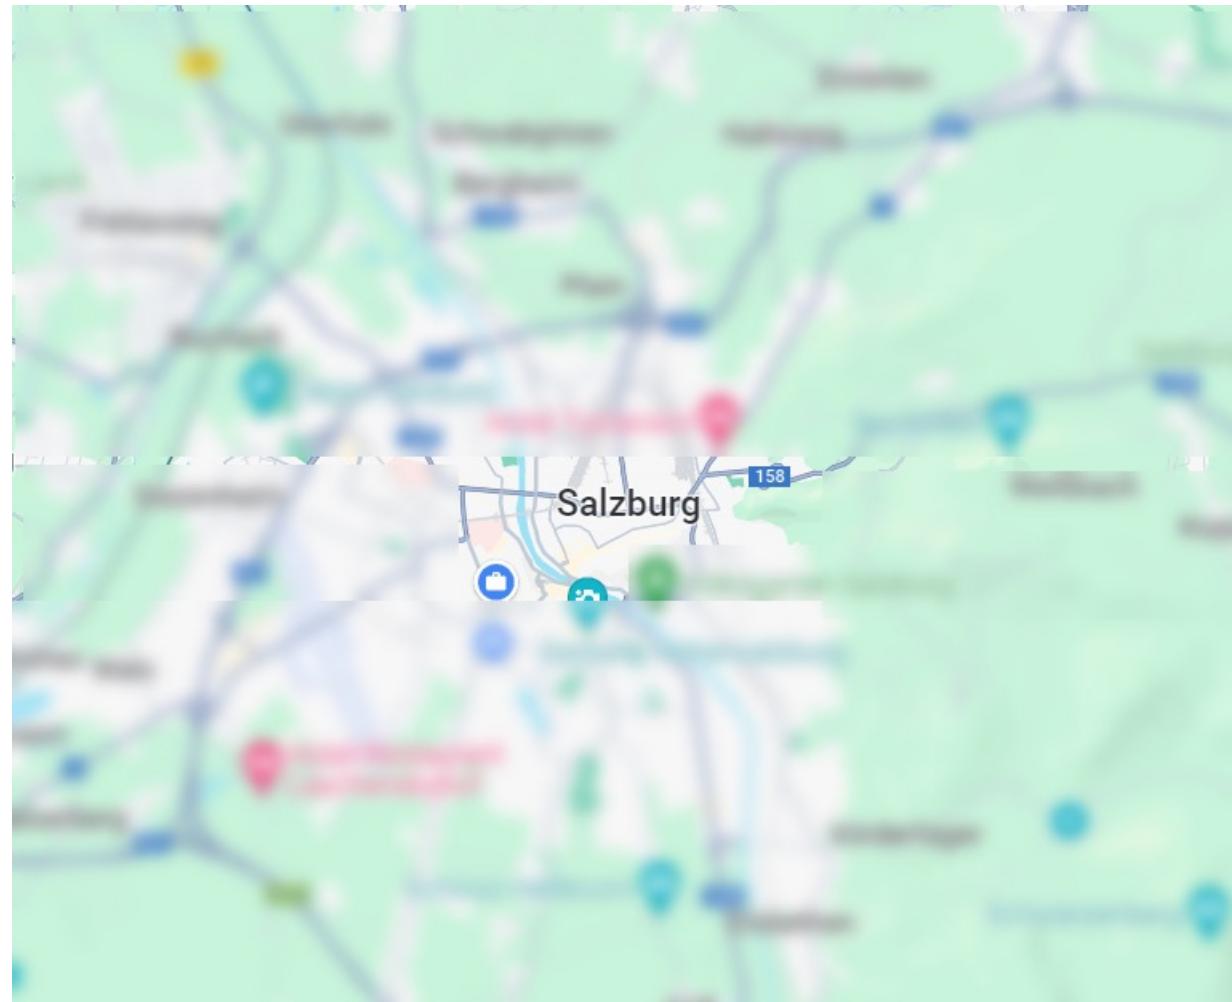


# Introducing SQL Server Query Optimizer

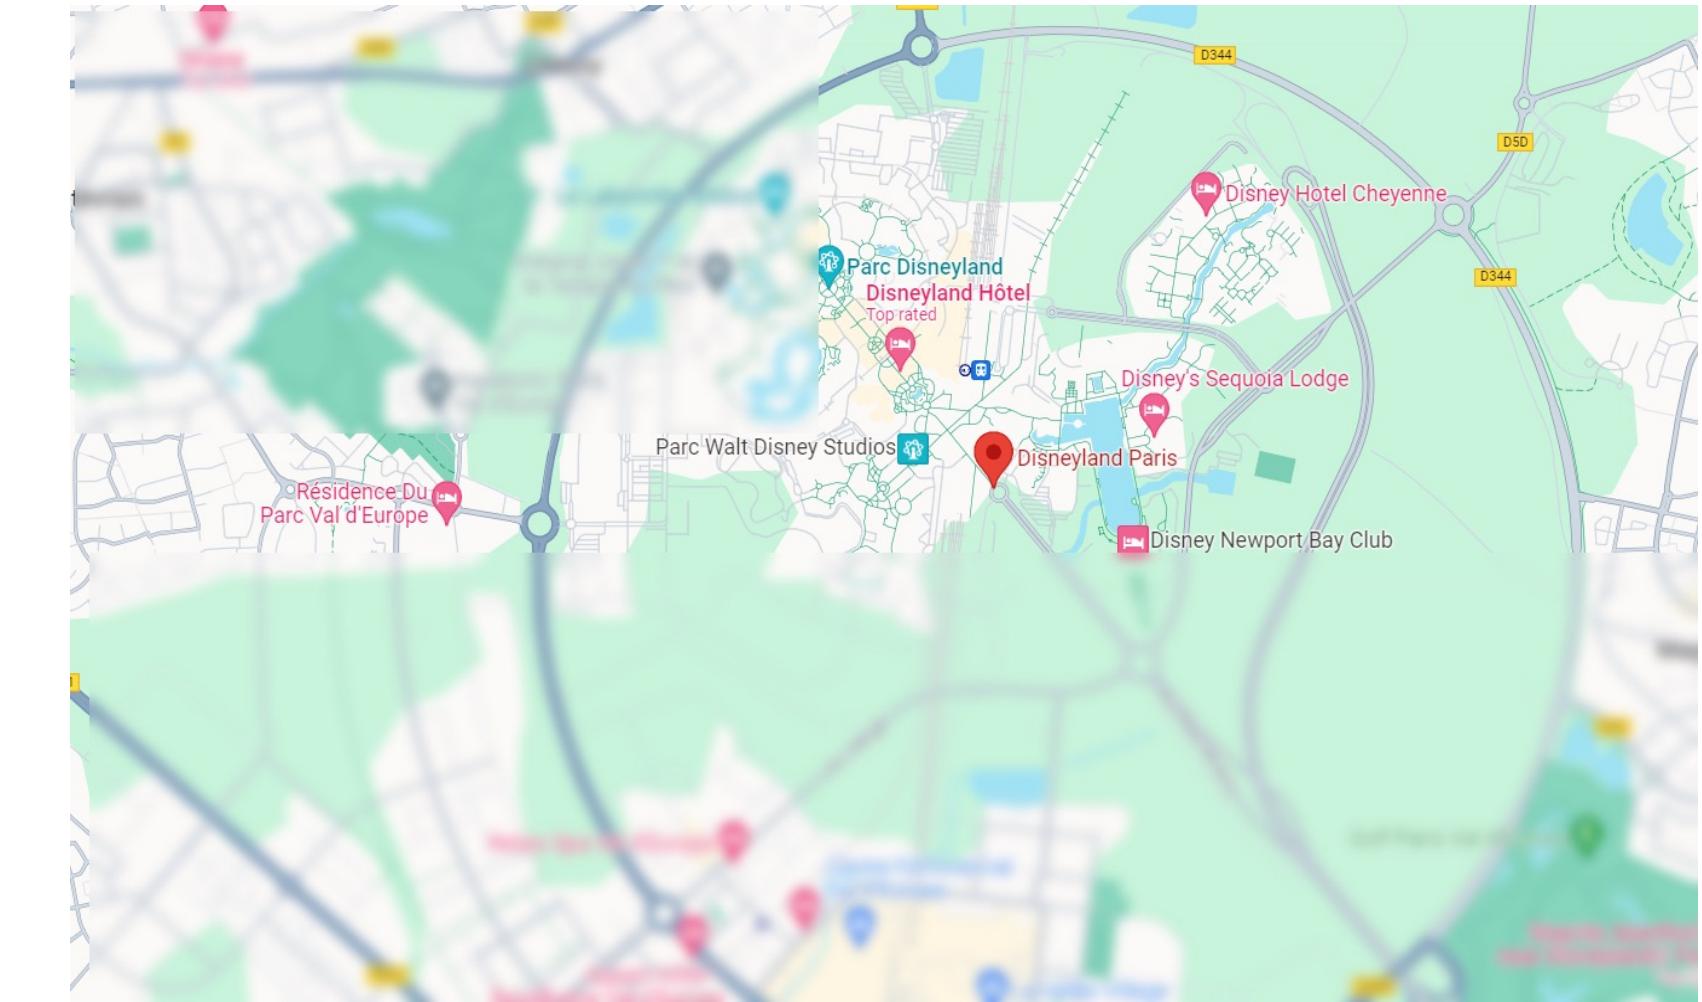


# Brain behind SQL Server query performance



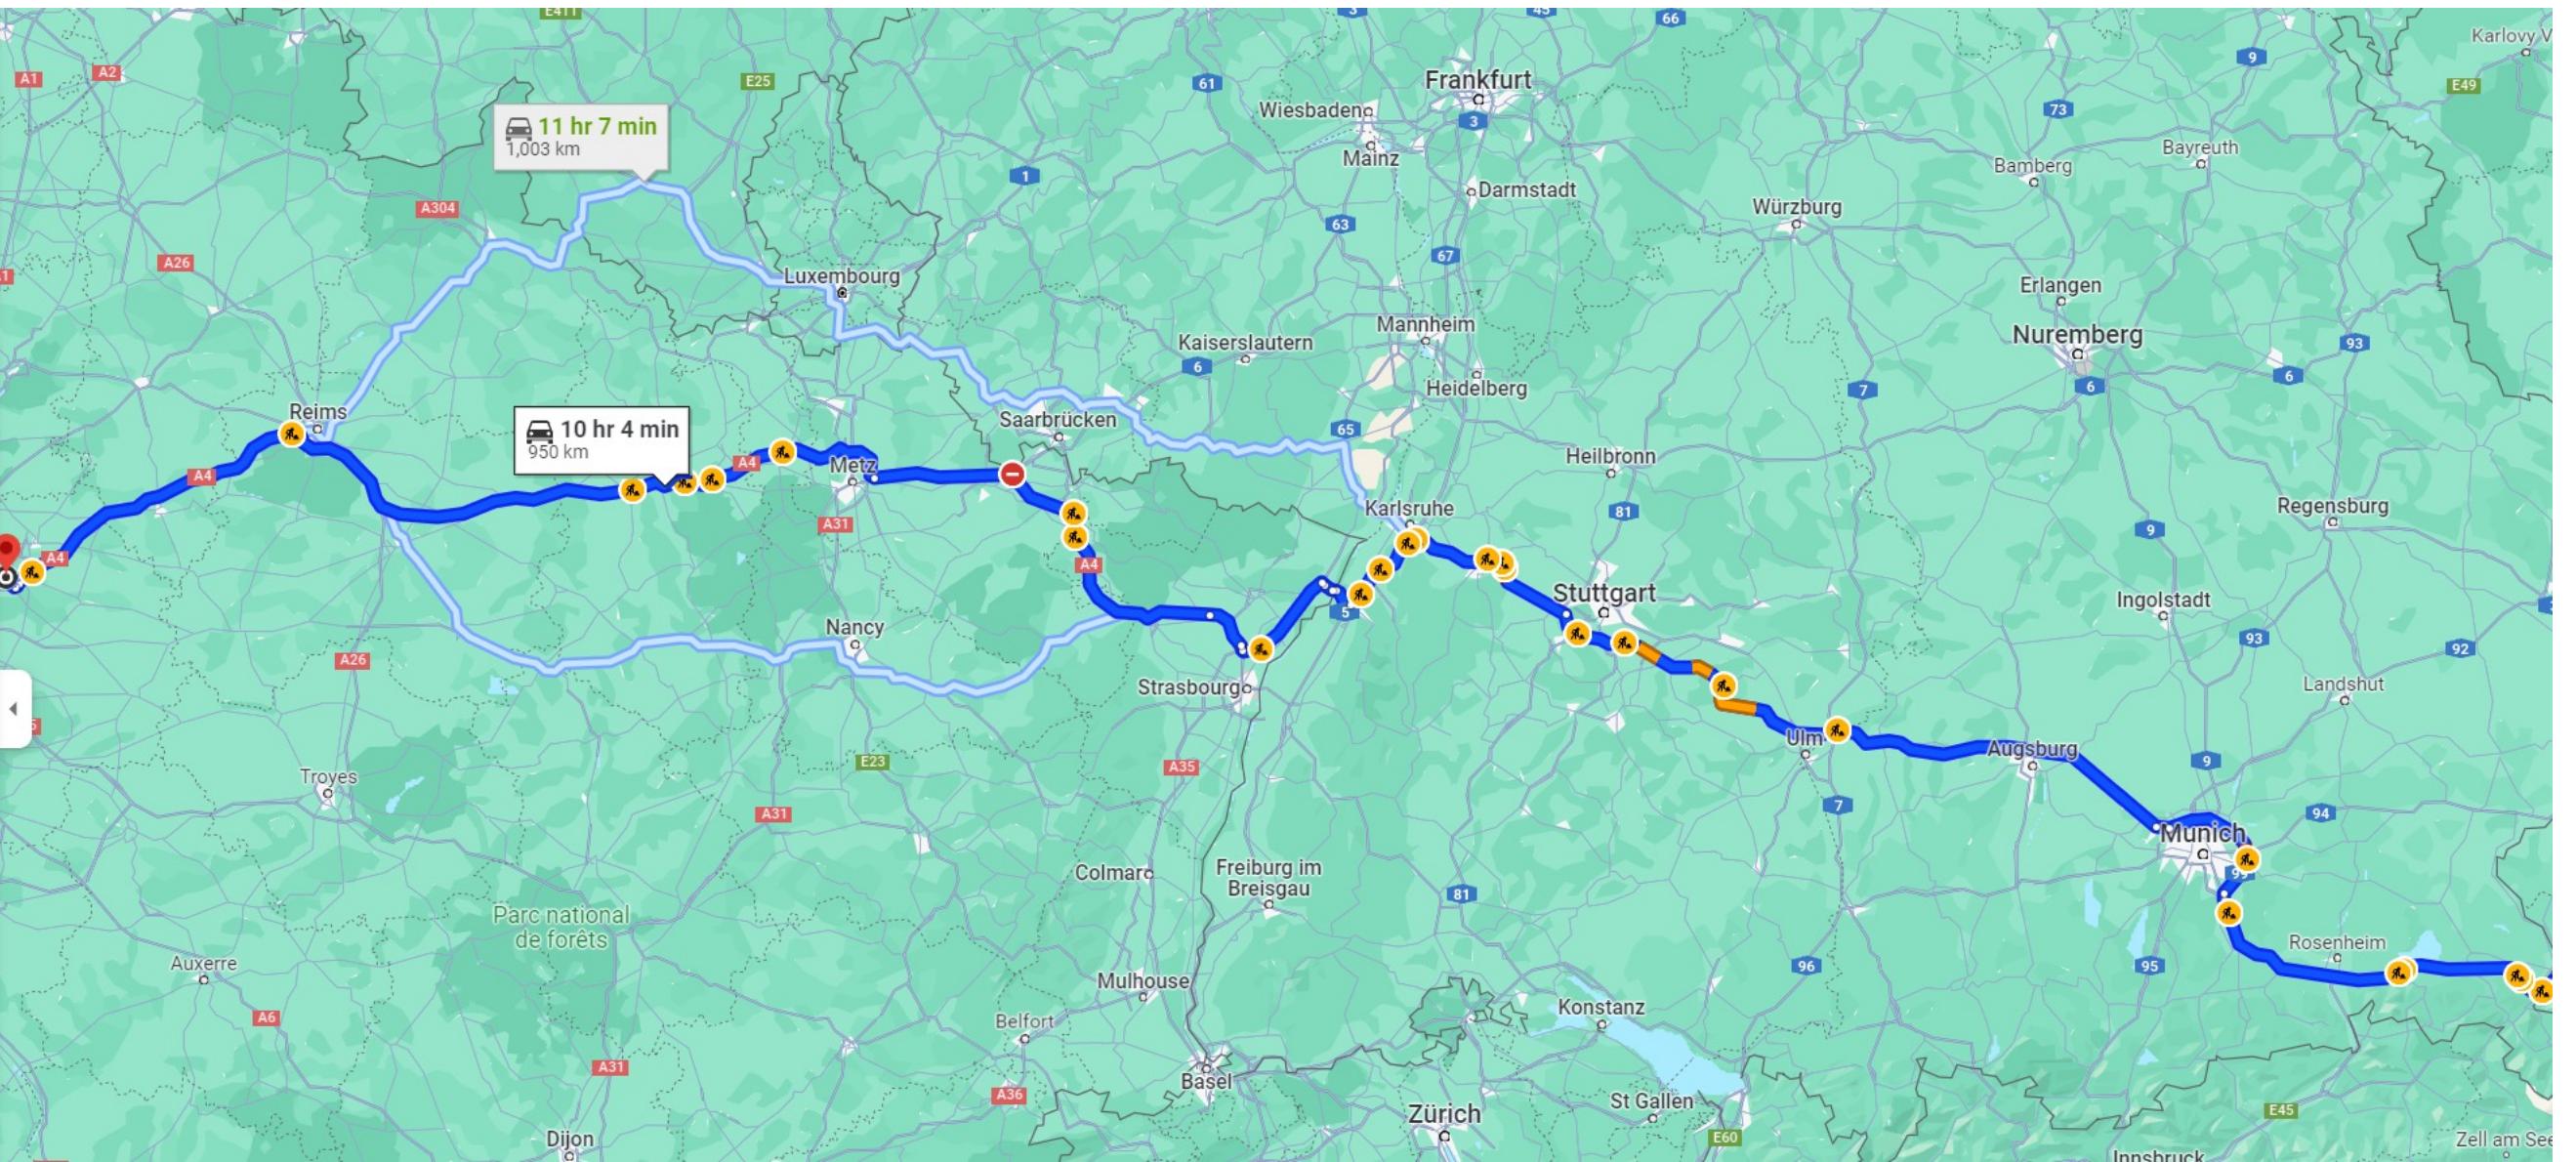


Salzburg, Austria



Disneyland, Paris





- ✓ **Distance**
- ✓ **Traffic jams**
- ✓ **Possible accidents**
- ✓ **Road works**
- ✓ **Road type**

- ✓ **Evaluate criteria**
- ✓ **Estimate the cost**



The “optimal” query plan is  
not always the most  
optimal!



# T-SQL Statement Evaluation Order

```
SELECT column1  
      , column2  
      , SUM(column3) AS sumCol  
  FROM table1 AS t1  
 INNER JOIN table2 AS t2 ON t1.id = t2.id  
 WHERE column1 = 'ABC'  
 GROUP BY column1  
      , column2  
 HAVING SUM(column3) >= 100  
 ORDER BY column2 DESC
```



# T-SQL Statement Evaluation Order

1

FROM

```
SELECT column1  
      , column2  
      , SUM(column3) AS sumCol  
FROM table1 AS t1  
INNER JOIN table2 AS t2 ON t1.id = t2.id  
WHERE column1 = 'ABC'  
GROUP BY column1  
      , column2  
HAVING SUM(column3) >= 100  
ORDER BY column2 DESC
```



# T-SQL Statement Evaluation Order

```
SELECT column1  
      , column2  
      , SUM(column3) AS sumCol  
  FROM table1 AS t1  
 INNER JOIN table2 AS t2 ON t1.id = t2.id  
WHERE column1 = 'ABC'  
 GROUP BY column1  
      , column2  
 HAVING SUM(column3) >= 100  
 ORDER BY column2 DESC
```

1 FROM

2 WHERE



# T-SQL Statement Evaluation Order

```
SELECT column1  
      , column2  
      , SUM(column3) AS sumCol  
  FROM table1 AS t1  
INNER JOIN table2 AS t2 ON t1.id = t2.id  
 WHERE column1 = 'ABC'  
GROUP BY column1  
      , column2  
 HAVING SUM(column3) >= 100  
 ORDER BY column2 DESC
```

- 1 FROM
- 2 WHERE
- 3 GROUP BY



# T-SQL Statement Evaluation Order

```
SELECT column1  
      , column2  
      , SUM(column3) AS sumCol  
  FROM table1 AS t1  
INNER JOIN table2 AS t2 ON t1.id = t2.id  
 WHERE column1 = 'ABC'  
 GROUP BY column1  
      , column2  
HAVING SUM(column3) >= 100  
 ORDER BY column2 DESC
```

- 1 FROM
- 2 WHERE
- 3 GROUP BY
- 4 HAVING



# T-SQL Statement Evaluation Order

```
SELECT column1  
      , column2  
      , SUM(column3) AS sumCol  
  FROM table1 AS t1  
 INNER JOIN table2 AS t2 ON t1.id = t2.id  
 WHERE column1 = 'ABC'  
 GROUP BY column1  
      , column2  
 HAVING SUM(column3) >= 100  
 ORDER BY column2 DESC
```

- 1 FROM
- 2 WHERE
- 3 GROUP BY
- 4 HAVING
- 5 SELECT



# T-SQL Statement Evaluation Order

```
SELECT column1  
      , column2  
      , SUM(column3) AS sumCol  
  FROM table1 AS t1  
INNER JOIN table2 AS t2 ON t1.id = t2.id  
 WHERE column1 = 'ABC'  
GROUP BY column1  
      , column2  
 HAVING SUM(column3) >= 100  
ORDER BY column2 DESC
```

- 1 FROM
- 2 WHERE
- 3 GROUP BY
- 4 HAVING
- 5 SELECT
- 6 ORDER BY



# **Understanding the Stages of the Query Optimization Process**

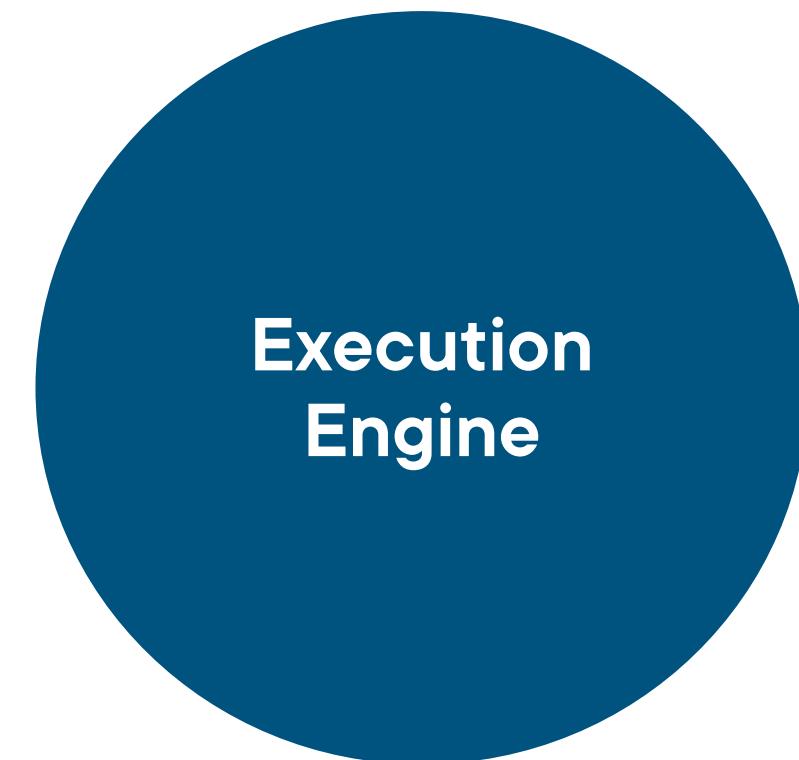


We only specify which data  
we want, not the **STEPS** to  
get this data



# Query Processor Engine in a Nutshell

- ✓ Accepts the query
- ✓ Generates the optimal query plan execution
- ✓ Executes the query plan



# Query Processing Stages

1

Parsing

~~SLECT~~

2

Binding

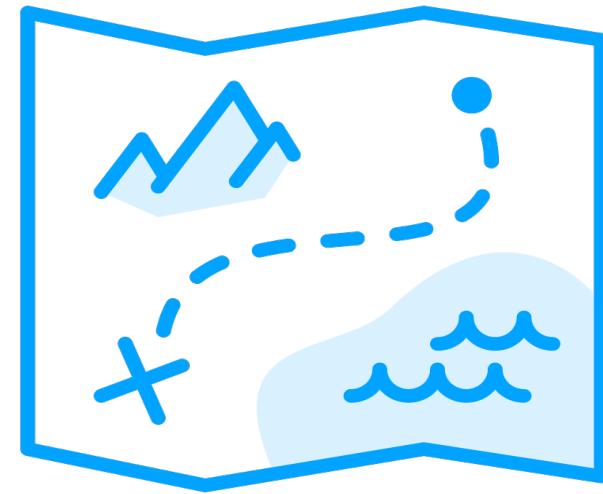
- ✓ Objects necessary for the query (tables, columns, indexes...)
- ✓ Identifying data types
- ✓ Processing aggregation types for GROUP BY

3

Optimization



# Optimization Stage Key Steps



## Generating possible execution plans

Number of candidate plans as a series of physical operations



## Cost evaluation for each plan

Plan with the lowest cost will be passed to the Execution Engine



# Query optimization

=

mapping logical query  
operations to physical  
operations



## Demo



### Capture query performance metrics

- Statistics
- QueryTimeStats
- Querying DMVs
- Query Store



# Summary



**Is 10ms fast or not?**

- Setting the right expectations

**The most common query bottlenecks**

**Similarity between Google Maps and SQL Server query plans**

- Get results using multiple different plans

**SQL Server Query Optimizer = cost based engine**

- Picks the optimal plan among many candidates



**Up Next:**

# **Understanding Query Plans**

---

